

# Third System — Cross-Category Findings

## AI Presence Index v0.6 across Five Categories

*Third System · 30 April 2026*

---

*For the first time, we can measure the moment that matters.*

**Five categories. Same patterns. AI mediation is producing structural effects on brand visibility that traditional measurement cannot capture, and the effects are now clear enough to act on.**

Third System has measured the AI Presence Index in five categories: project management software, running shoes, premium olive oil, premium facial skincare, and personal finance applications. The categories were chosen to span the broadest possible range of properties — software and physical goods, B2B and consumer, mature and emerging, identity-light and identity-heavy, US-centric and globally fragmented, low regulatory salience and high. The five differ on every meaningful axis except one: each is a category in which a meaningful share of consumers now consult AI before deciding what to buy.

Across all five, six patterns hold. The first five replicate strongly, in some cases across all five categories. The sixth emerged from a single category but is the most consequential single finding in the dataset.

First, AI Presence varies meaningfully across models for the same brand, and the size of the variance follows the coherence of the category's discourse. Personal finance, the most fragmented discourse in the dataset, produced a 71-percentage-point spread for Rocket Money — the largest single-brand variance observed across any category. Skincare, the most coherent discourse (dermatology consensus, ingredient-focused, science-positioned), produced single-digit spreads for top brands. Olive oil and project management software, both fragmented in different ways, produced spreads of 40 and 41 points. Running shoes, with its converged technical-review consensus, produced narrow spreads. The relationship between discourse coherence and per-model variance is now strongly supported across all five categories.

Second, AI converges on incumbents when forced to compare and elevates challengers when asked about emerging tools. The asymmetry repeats across all five categories. The strength varies — sharpest in PM software, running shoes, and olive oil; weaker in skin-care where no clear “emerging” tier exists. Across the five categories, brands that surface in Discovery contexts are systematically different from brands that surface in Comparison contexts, often dramatically so.

Third, AI brand visibility diverges from consumer awareness, and the gap intensifies with category age and discourse maturity. Nike places sixth in running shoes despite the largest

marketing budget in the category. Bertolli, Colavita, and Goya combine for less than 8% AI Presence in olive oil despite dominating supermarket distribution. La Mer, SK-II, and the entire luxury heritage skincare tier combined are outscored ten-to-one by CeraVe alone. In personal finance, the divergence inverts in a different way: YNAB, with roughly one million paid subscribers, achieves 100% AI Presence, while Rocket Money, with five times more users, surfaces inconsistently across models.

Fourth, Default Reinforcement intensifies when AI training data carries strong consensus aligned with the prompt's framing. In skincare, when asked what dermatologists recommend, three drugstore brands hit 100% mention rate — the AI's response was nearly deterministic. In personal finance, when asked what people good with money use, YNAB and Empower hit 100%. The consensus brands win the identity prompts so completely that brands outside the consensus simply don't appear, regardless of marketing weight.

Fifth, AI brand visibility appears to carry a discourse-language bias. Spanish olive oil brands underperformed substantially in a category Spain dominates globally. Korean skincare brands disappeared entirely in a category where K-beauty is well-covered in English. The pattern is preliminary — only two categories produced this signal — but the consistency across two categories with different language profiles strengthens the finding from where it stood at one category. A category designed specifically to test the hypothesis is planned for Phase 2 continuation.

Sixth, and most consequentially, AI training data lag creates phantom brand presence. Mint, the personal finance application shut down by Intuit in March 2024, was mentioned in 44% of AI responses about personal finance in measurements taken two years after its decommissioning. On Anthropic specifically, Mint surfaced in 65% of responses — a brand that no longer exists, mentioned in two-thirds of recommendations about live products. The implication is direct: AI brand visibility is its own time-lagged surface, and brands that have shut down, rebranded, or merged can persist for years in the recommendation tier. Marketers may be optimizing against competitors that no longer exist, while their own past marketing investment continues to generate AI presence after their products are discontinued. In AI mediation, brands die slowly.

The implication for brand strategy compounds across these six findings. Marketers operating with brand-tracking instruments built for the human-mediated era are measuring a smaller and time-shifted surface than the one their consumers actually navigate. Brands that lead in awareness can trail in AI recommendation. Brands that lead in AI can be unknown to broader markets. Brands that have ceased to exist can still dominate AI mediation. Every brand valuation model in use today was built for a world where humans browsed. The AI Presence Index is the adjustment for a world where they ask.

Third System publishes all five category measurements alongside this summary. Methodology, full leaderboards, and the underlying datasets are available at [thirdsystem.ai](https://thirdsystem.ai). Phase 3 of the program will correlate these measurements against external brand-tracking data to test whether AI Presence functions as a leading indicator of consumer behavior or whether it measures AI behavior in isolation. Both findings would be useful, and neither has been established.

## **What we measured**

This report combines findings from five Third System category measurements completed in late April 2026.

The five categories were chosen sequentially, each selected to test framework properties not represented by prior measurements. Project management software was the inaugural category, selected for a mature B2B competitive set with high AI delegation by buyers. Running shoes followed, selected for a contrast on physical-versus-software, consumer-versus-B2B, and substantially higher identity load. Premium olive oil was added third, selected to introduce a category in which English-language editorial discourse dominates a globally distributed production base. Premium facial skincare was added fourth, selected to test a category with extremely coherent editorial consensus dominated by dermatological recommendation logic. Personal finance applications were added fifth, selected to test a category with regulatory salience, fragmented discourse, and recent market disruption.

The five categories collectively span every meaningful axis the framework claims to apply across: identity load (low to very high), buyer type (B2B to consumer), domain (software, durables, food, beauty, financial services), discourse coherence (very low to very high), production geography (US-centric to globally fragmented), and regulatory salience (low to high).

All five measurements followed the same methodology. Six prompts mapped to Category Entry Points were issued eight times each to two frontier AI models, OpenAI's gpt-5.4-mini and Anthropic's claude-sonnet-4-6, at temperature 0.7. Brand mentions were extracted via structured-output classification and mapped against a category-specific brand registry. Project management software covered 19 brands, running shoes 17, premium olive oil 20 (after one registry revision), premium facial skincare 31 (after one registry revision), and personal finance applications 16. Each category produced 96 successful measurements. The headline metric in each case is Presence: the raw mention rate of each brand across all measurements, reported on an absolute 0-100 scale without normalization.

Two registry revisions occurred during the measurement period. In premium olive oil, the first run surfaced Cobram Estate as the most-mentioned brand outside the registry; the registry was expanded to include it and three additional brands present in the unknown-mentions list. In premium facial skincare, the first run surfaced La Roche-Posay, Vani-cream, and six additional dermatologist-recommended drugstore brands as significant unknowns; the registry was expanded to include all of them. The published data for both categories comes from the post-revision runs. Earlier runs are preserved in the methodology log for traceability. The unknown-mentions list serves as a continuing diagnostic for registry quality, and registry expansion is now a standard practice across measurement cycles.

The full methodology log, including version history and the documented v0.2 calibration error, is published as `methodology.md` v0.3. All five category leaderboards and the underlying response datasets are available alongside this report.

Five categories is the threshold at which cross-category claims can move from suggestive to substantively defensible. The findings reported here cross that threshold. They are not yet at the threshold of universal generalizability, which would require both broader category coverage and longitudinal data across model updates. Phase 2 continuation will introduce both.

### Three modes of AI response

When we measure AI Presence, we are not measuring a single behavior. Across the five categories Third System has measured, the AI engages three structurally distinct modes when responding to category prompts, and which mode it engages depends on the fram-

ing of the question rather than the structure of the category. Recognizing these modes is foundational to interpreting everything that follows in this report.

**Brand mode** is what most of this report measures. The AI is asked about a category, and it names brands within that category. CeraVe at 100% in skincare's Identity prompt, YNAB at 100% in personal finance's Identity prompt, Linear and Asana at 100% in project management software's Comparison prompt — all of this is brand mode. The AI treats the category as a marketplace of named producers and answers accordingly.

**Component mode** appears when the prompt activates a clinical, technical, or specifications frame. The AI shifts from naming brands to naming ingredients, materials, or product properties. In skincare, the Contextual prompt about fine lines in one's late thirties produced this mode 95% of the time. Across both Anthropic and OpenAI, the AI named SPF 30+, retinol, tretinoin, vitamin C, niacinamide, hyaluronic acid, and peptides. The AI named almost no brands. It treated the question as a dermatological consultation, in which the answer is *what's in the bottle*, not *whose bottle to buy*. Olive oil showed component mode at lower intensity in the Functional prompt (62%): the AI named "extra virgin olive oil," "EVOO," "light olive oil," and "refined olive oil" instead of producer brands. The Comparison prompt for olive oil showed component mode in 62% of empty-canonical responses, with the AI distinguishing "single-origin," "estate-bottled," and "PDO/PGI" categories rather than naming specific estates.

**Authority mode** appears when the prompt asks about discovery, novelty, or where to learn. The AI shifts from naming brands to naming publications, communities, retailers, and competitions. In skincare's Discovery prompt, the AI exclusively named *Allure*, *Byrdie*, *r/SkincareAddiction*, *Reddit*, *WWD Beauty*, *Vogue*, *Sephora*, and *Ulta*. In olive oil's Discovery prompt, the AI exclusively named *Olive Oil Times*, *NYIOOC World Olive Oil Competition*, *Flos Olei*, *Bon Appétit*, *Saveur*, and *Zingerman's*. Asked where the new and noteworthy can be found, the AI named the discourse infrastructure rather than the products. It treated the question as one of orientation toward a category's information sources, not as one of product recommendation.

The implication for measurement is structural. A leaderboard cell with zero brand mentions does not necessarily indicate a brand's absence from the AI's awareness. It can indicate that the prompt activated a different mode entirely, in which the AI was making a different kind of recommendation. Component mode and authority mode are not measurement failures. They are the AI engaging the question on a different layer of category understanding, and the framework's job is to recognize which layer is active.

The implication for brand strategy is broader. AI brand visibility is one of three surfaces on which a brand competes in the AI tier. The second surface is ingredient or component association: the brands a consumer associates with retinoids will benefit from any prompt that activates component mode in skincare, regardless of the brands' overall AI Presence. The third surface is authority and publication relationships: the brands featured by the publications the AI names in authority mode have a privileged path into the AI's discovery recommendations. A brand strategist who optimizes only for the first surface — direct brand mentions — will under-invest in the other two.

Future Phase 2 categories will introduce a mode classifier as a pre-step in the measurement process, allowing each prompt's response to be scored on which mode it activated, with brand-mode Presence reported alongside the rate of component-mode and authority-mode activation. This refinement strengthens the framework's epistemic architecture without requiring re-measurement of existing categories: the modes are visible in the data already

collected, and the existing Presence scores remain valid as brand-mode-conditional measurements.

What this section establishes will reappear in the patterns that follow. Pattern 4 (discourse-language bias) is strengthened by the observation that authority-mode responses in two unrelated categories named exclusively English-language publications. Pattern 5 (Default Reinforcement) is sharpened by the recognition that the AI's "default" can be a mode rather than a brand. Pattern 1 (per-model variance) reads differently when one accepts that two AIs may be in different modes for the same prompt rather than disagreeing within brand mode. The three-mode structure is the foundation; the patterns are observations across that foundation.

### **Pattern 1: AI Presence varies meaningfully across models, in proportion to discourse fragmentation**

Across all five categories, the same brand can score significantly higher on one AI than another. The variance is consistent enough to be measured, large enough to matter for strategy, and follows a predictable relationship with the structure of the underlying discourse.

The largest single-brand variance observed across the five-category dataset comes from personal finance. Rocket Money, a subscription-management and budgeting application owned by Rocket Companies, was mentioned in 71% of OpenAI responses about personal finance applications and zero percent of Anthropic responses. Two AI models, asked the same six prompts about the same category, produced effectively disjoint answers about whether Rocket Money exists as a category player. Mint, the defunct Intuit application, showed a 42-point spread in the same direction (Anthropic 65%, OpenAI 23%). Five other personal finance brands had spreads above 14 points.

Premium olive oil produced spreads nearly as wide. Frescobaldi Laudemio, a Tuscan estate brand with strong editorial presence and limited US retail distribution, appeared in 42% of Anthropic responses and 2% of OpenAI responses, a forty-point spread. Castillo de Canena, a Spanish premium estate brand, showed the inverse asymmetry. Frantoio Muraglia, an Italian premium brand, appeared in 15% of Anthropic responses and 42% of OpenAI responses.

Project management software showed similar magnitudes for top brands. ClickUp surfaced in 83% of OpenAI responses and 42% of Anthropic responses, a forty-one-point spread. Asana showed inverse asymmetry, dominant on OpenAI at 94% and softer on Anthropic at 67%.

Running shoes and skincare produced markedly narrower variance. The largest top-brand spread in running shoes was Hoka at 19 points. In skincare, CeraVe registered identical 67% Presence on both models — zero spread, the only zero-spread leader observed across any category. The largest top-brand spread in skincare was Paula's Choice at 23 points.

The pattern across five categories now points clearly to a structural relationship. Variance is not a function of category maturity, the v0.4 hypothesis: olive oil is mature and produces high variance, while running shoes is similarly mature and produces low variance. The unifying variable is **discourse coherence**. Personal finance has the most fragmented discourse in the dataset — split across YouTube creators, traditional financial media, fintech press, Reddit communities, and influencer recommendations, none of which agree about which brands deserve consideration. Olive oil discourse fragments across geographies (Italy, Spain, California, Australia), production traditions, and audience tiers. Project

management software discourse fragments across competing professional frames. Running shoes discourse converges around a small number of authoritative voices speaking in one language about a small set of technical criteria. Skincare discourse converges even more tightly around dermatological recommendation logic. The two converged-discourse categories produce narrow variance; the three fragmented-discourse categories produce wide variance.

The relationship is now strongly supported across all five categories. The implication for brand strategy is that per-model variance is not a transient artifact of any specific category's youth, but a structural property of how the AI tier processes the underlying discourse. Brands operating in fragmented-discourse categories should expect persistent variance across models and plan AI-channel strategy with that variability in mind. Brands in converged-discourse categories will see narrower variance but should not mistake convergence for stability — model updates can shift consensus quickly, and convergence makes those shifts more visible.

A practical consequence holds in every category measured. A marketer running an AI-visibility audit on a single model is measuring a slice of a larger and more variable surface. Multi-model measurement is not a refinement of the methodology. It is the methodology.

## **Pattern 2: The AI's answer depends on the question**

Across all five categories, the AI converges on incumbent brands when forced to compare options and elevates a sharply different set when asked about emerging tools. The asymmetry is consistent. The strength varies.

In premium olive oil, the comparison prompt returned California Olive Ranch and Cobram Estate at near-saturation; the discovery prompt elevated Brightland, Graza, and Kosterina, three direct-to-consumer brands whose combined sales are a fraction of the comparison-prompt leaders. In running shoes, the comparison prompt returned Nike, Asics, and New Balance at 100%; the discovery prompt returned Norda, a four-year-old Canadian trail-running brand at 100%. In project management software, comparison returned Asana, Monday, and Jira at 100%; discovery surfaced Linear and Height. In personal finance, comparison returned YNAB, Empower, and Monarch Money; discovery (which named Mint by name) elevated Monarch and Copilot.

Skincare is the structural exception. The comparison prompt produced strong convergence on CeraVe, Drunk Elephant, and Skinceuticals. But the discovery prompt did not produce a clear “emerging tier” — Tatcha, Augustinus Bader, and Youth to the People surfaced at 25%, 25%, and 12% respectively, with no brand achieving the 50% threshold seen in other categories. The asymmetry exists in skincare; it is just weaker. The mechanism appears to be that skincare's editorial discourse has been so thoroughly worked through that no brand confidently occupies an “emerging” cognitive slot in the AI's recommendation logic. Every plausible “emerging” skincare brand has been editorially established for half a decade or more.

Three findings sit inside this asymmetry, all of them strengthened by five-category replication.

The first is that the AI's working definition of “emerging” is unstable across categories and consistently does not mean “new.” In project management software, the AI's emerging-tools list included Notion, valued in the tens of billions of dollars. In olive oil, the discovery-prompt leaders Brightland and Kosterina are seven and eight years old respectively. In per-

sonal finance, the discovery prompt named YNAB at 100% — a 17-year-old company. The AI's working definition of “emerging” appears to mean “non-dominant in the comparison frame,” not “young,” “small,” or “novel.”

The second is that challenger brands can achieve concentrated AI presence in narrow Category Entry Points without registering at all on broader prompts. Norda's overall Presence in running shoes is 17%; its Discovery Presence is 100%. Kosterina's overall Presence in olive oil is 27%; its Discovery Presence is 60%. Augustinus Bader's overall Presence in skincare is 7%; its Discovery Presence is 25%. The leverage available to a challenger brand in a single CEP can be much larger than its overall AI visibility would suggest, and that leverage is the right strategic target.

The third is that prompt design materially affects which version of the asymmetry shows up. Discovery prompts that name a specific market disruption (the personal finance prompt naming Mint by name) produce strong convergence. Discovery prompts that ask generically about “emerging brands” produce diffuse responses, particularly in categories with no clearly defined emerging tier. Brand strategy that tests AI presence with generic discovery prompts will under-detect the visibility available in narrow CEPs that more concrete framings would surface.

The strategic implication for any challenger brand follows from these three findings. Pursuing AI visibility as a category-level objective will under-serve any brand that is not already an incumbent. The visibility worth pursuing is contextual, varies by the cognitive frame of the consumer query, and is most accessible to challengers in narrow CEPs where the AI's frame is more permissive. Strategists who treat AI presence as a single number will misread the surface they are operating on.

### **Pattern 3: AI brand visibility diverges from consumer awareness, in two directions**

Across all five categories, AI brand visibility does not track consumer awareness, and the gap takes different forms depending on the category's structure. In four of the five categories, the divergence runs one direction: brands with high consumer awareness under-perform their market position in AI mediation. In personal finance, the divergence inverts: a brand with modest market share dominates AI mediation completely. The framework predicts both directions.

The four downward-divergence cases form a clear pattern. Premium olive oil shows it in extreme form: the four most distributed legacy supermarket brands — Bertolli, Colavita, Goya, and Filippo Berio — combine for less than 8% AI Presence in aggregate. Goya, despite being one of the largest Latin foods brands in the United States, scored exactly zero. Brightland, a seven-year-old direct-to-consumer brand that exists in only a small fraction of those brands' retail footprint, scored 46%. AI mediation prefers the seven-year-old.

Skincare shows it in even sharper form. La Mer, SK-II, Lancôme, Estée Lauder, Clinique, and Olay — six legacy luxury heritage brands — combine to roughly 6% Presence. CeraVe, the dermatologist-recommended drugstore brand, scored 67% alone. The legacy luxury tier in skincare is outscored ten-to-one by a single drugstore brand. La Mer's \$400 cream loses to CeraVe's \$15 cream at thirty-to-one in AI presence ratio.

Running shoes shows it in less extreme but still strong form. Nike has the largest marketing budget of any brand in the running-shoe category and arguably the highest cultural saturation of any brand on Earth. In AI Presence, Nike places sixth at 57%, behind Asics, Brooks, Saucony, New Balance, and Hoka. Brooks, a brand whose marketing budget is a fraction of

Nike's, ties for first at 80%.

Project management software shows the divergence in mild form, consistent with v0.4. AI presence and traditional consideration tracked more closely at the top of the leaderboard. The asymmetry is not absent — Linear's AI Presence runs ahead of its current market share — but the gap is narrow.

The fifth case, personal finance, inverts the direction. YNAB, with approximately one million paid subscribers, scored 100% AI Presence across every prompt and every model. Rocket Money, with five times more users (over five million), surfaced in zero percent of Anthropic responses and 71% of OpenAI responses, depending on the model. EveryDollar, the budgeting application from Ramsey Solutions with millions of users in its ecosystem, scored 25%. Empower, the platform for higher net-worth users, scored 66%. The brand with the smallest user base in the top tier dominates the category in AI mediation; brands with substantially larger user bases trail or splinter across models.

The mechanism connecting these five cases is consistent. AI training data accumulates editorial framings over time. Where editorial discourse builds a coherent quality narrative aligned with a specific brand's positioning, that brand's AI presence outpaces its commercial scale. CeraVe is not the largest skincare brand by revenue; it is the most consistently recommended brand in dermatological discourse. YNAB is not the largest personal finance brand by users; it is the most methodologically distinctive brand in the discourse about financial discipline. The AI tier privileges discourse position over distribution scale.

The pattern holds in both directions. **In categories where editorial discourse has built up coherent quality narratives, AI mediation rewards discourse positioning regardless of market share — upward when small brands have strong narratives, downward when large brands lack them.** This is the same finding pointing two directions, and it now holds across all five categories.

The implication for brand strategy is that AI visibility is a function of discourse position, not commercial position, and the two are not necessarily aligned. In some categories they correlate weakly. In others they correlate inversely. The CMO of a brand whose marketing investment optimizes for distribution scale is not necessarily building AI presence. The CMO of a brand whose investment optimizes for editorial positioning may be building substantial AI presence on a small commercial base. The instruments of brand strategy need to read both surfaces, and the reading needs to be specific to the category being measured.

#### **Pattern 4: AI brand visibility appears to carry a discourse-language bias**

The fourth pattern emerged in two of the five categories with consistent direction, in different national-origin contexts. It remains preliminary because no measurement to date has been designed specifically to test the hypothesis. The consistency of the signal across two unrelated categories strengthens it from where it stood at one category, while leaving room for refinement.

In premium olive oil, three brands were added to the registry specifically to ensure Spanish representation. Spain produces more olive oil than any other country in the world, accounting for roughly 45% of global production. Goya is one of the most widely distributed Spanish-affiliated brands in the United States. Castillo de Canena is a premium Andalusian estate brand with growing US specialty distribution. Núñez de Prado is a respected family-owned Andalusian producer with strong editorial presence in Spanish-language food media. Their AI Presence scores were 0%, 23%, and 5% respectively. None of the three placed

in the top five. The category leaders were California Olive Ranch (US), Cobram Estate (Australia), Brightland (US), Frantoio Muraglia (Italy), and Graza (US). Italian brands collectively performed in line with their commercial profile in the US market. American and Australian brands over-performed relative to global production share. Spanish brands under-performed substantially.

In premium facial skincare, Beauty of Joseon was included in the registry as a representative of K-beauty's strongest current brand presence in US specialty retail. Korean cosmetics enjoy significant English-language coverage in major beauty media — Allure, Vogue, Glamour, and The Strategist have published extensively on K-beauty for over a decade. Beauty of Joseon scored zero across both runs of the skincare measurement. The skincare unknowns list, which surfaced eight additional missing brands, contained no non-English-discourse brands at all. Every brand the AI surfaced as a category player was an American DTC brand or a French/American clinical brand.

The most defensible explanation for both results is that AI training corpora over-represent English-language discourse on these categories, and the bias persists even when the cultural reference of the brand is well-covered in English. American food magazines, Italian food media translated into English, and US-based recipe websites dominate the digital corpus on olive oil. American beauty magazines and US-based dermatology blogs dominate the corpus on skincare. Spanish-language food media and Korean-language beauty media, even when authoritative within their domestic markets, do not propagate into English-language AI training data at the rates that comparable English-language sources do. The AI's recommendations therefore reflect the geography and language of the discourse it was trained on, not the geography of the production it is asked about.

A subtle refinement of the hypothesis becomes visible across the two cases. Tatcha, a Japanese-themed skincare brand founded in San Francisco with American-English marketing, scored 18% Presence in skincare. Beauty of Joseon, a Korean brand with Korean-language primary marketing despite some US-targeted English content, scored zero. The cultural reference (Japanese aesthetic) does not seem to be the constraint. The marketing-discourse language does. **The hypothesis, refined: AI under-surfaces brands whose primary marketing discourse is conducted in a language under-represented in AI training data, even when the cultural reference is well-known in English.**

The audit of authority-mode responses across the two affected categories provides additional, unexpected evidence for the discourse-language hypothesis. When the AI was asked about discovery in skincare and olive oil, it named publications, communities, and retailers rather than brands. Every publication the AI named was English-language: *Allure*, *Byrdie*, *Vogue*, and *WWD Beauty* in skincare; *Olive Oil Times*, *Flos Olei*, *Bon Appétit*, and *Saveur* in olive oil. Spanish-language food media did not appear. Korean-language beauty media did not appear. Japanese, Italian, French, and German publications respected within their domestic categories did not appear. The AI's recommended path to category orientation runs entirely through English-language sources, even in categories where authoritative non-English sources demonstrably exist. This is not a finding about which brands the AI surfaces; it is a finding about which infrastructure the AI treats as legitimate. The discourse-language bias in brand recommendations is consistent with a deeper bias in which media systems are recognized as authoritative.

This finding remains preliminary in three ways. It is based on two categories with single nationally-skewed signals each. It involves small absolute numbers of test brands. And the alternative explanations — US retail availability, marketing investment in English-targeted

channels, or category-specific distribution structures — have not been controlled for. A more rigorous test would require measuring a category in which a non-English-discourse country dominates global production but has limited US-targeted English-language marketing, and observing whether the same asymmetry persists. Japanese kitchen knives, French wine, or Korean small electronics would each provide such a test. Phase 2 continuation will include at least one such category.

If subsequent measurements confirm the pattern, the implications are wide. AI mediation may systematically under-surface brands whose primary discourse is conducted in languages or media systems under-represented in AI training corpora. For global brands, this would mean that domestic-language brand-building does not propagate into international AI mediation at the same rate that English-language brand-building does. Brand strategy that ignores this asymmetry will under-invest in English-language presence relative to its strategic value, particularly for brands whose AI-mediated visibility matters in markets outside their primary linguistic context.

This is hypothesis at the edge of finding. Two categories support it. A third designed-for-test category will determine its status.

#### **Pattern 5: Default Reinforcement intensifies with discourse-prompt alignment**

The fifth pattern is closely connected to Pattern 1 (discourse coherence) and Pattern 3 (AI vs awareness divergence) but distinct from both. It describes how the strength of incumbent dominance in AI mediation depends not just on the underlying discourse coherence, but on how well-aligned that discourse is with the framing of the prompt itself.

Skincare provides the cleanest illustration. The Identity prompt asked: “What skincare brands do dermatologists actually recommend to their patients?” CeraVe, Cetaphil, and Neutrogena each appeared at 100% mention rate. Three brands tied at perfect saturation in 96 measurements. The Constraint prompt asked about sensitive skin: CeraVe and La Roche-Posay each appeared at 100%, with Vanicream at the same level. The Functional prompt asked about moisturizers for normal-to-dry skin: CeraVe at 100%. Across three different prompts framed in three different ways, the answers converged on the same small set of dermatologist-recommended drugstore brands.

The mechanism is visible: skincare’s editorial discourse has built a tight, well-documented consensus that dermatologists recommend specific clinical-positioning brands, and the prompts in this study repeatedly aligned with that consensus. When the AI’s training data carries strong consensus and the prompt’s cognitive frame activates that consensus, the AI’s response is nearly deterministic. There is no observable randomness; the same brands appear at saturation across measurements.

Personal finance shows a parallel pattern with different content. The Identity prompt asked about apps that “people who are actually good with money use,” loaded but ambiguous identity language. YNAB and Empower both scored 100% on this prompt. The Comparison prompt produced YNAB, Empower, and Monarch Money each at 100%. The Constraint prompt about partners managing money together produced YNAB and Monarch Money each at 100%. Across multiple identity frames, the AI converged on the same small brand set with deterministic confidence.

The contrast with running shoes and PM software is informative. In running shoes, the Identity prompt (“what serious competitive runners actually wear”) produced Nike at 100% — but no other brand reached 100% on that prompt. The category’s discourse is more

contested between cultural-narrative (“Nike is what serious runners wear”) and technical-recommendation (“Asics and Brooks are what serious runners wear”) framings. AI surfaces both, depending on which CEP is activated. The discourse exists; the prompt-discourse alignment is weaker.

In PM software, the Identity prompt produced multiple brands across both models without any single brand reaching saturation. The category’s discourse fragments across multiple identity-loaded frames (developer-first, marketing-first, generalist), none of which dominates. Default Reinforcement is observable, but compressed.

The cross-category data now suggests a structural relationship: **the strength of Default Reinforcement at the prompt level is a function of how completely the category’s editorial consensus aligns with the prompt’s framing.** When alignment is high (skincare’s dermatologist prompts, personal finance’s discipline prompts), AI responses are nearly deterministic. When alignment is partial (running shoes’ identity prompt, PM software’s frame-divided prompts), AI responses are stratified. When alignment is low (any prompt that activates a discourse the AI’s training data does not strongly carry), AI responses are diffuse.

The strategic implication for brand strategy is consequential and practical. A brand attempting to build AI presence cannot do so by general marketing investment; it must build editorial position aligned with the specific cognitive frames consumers are likely to use when prompting AI in the category. Generic awareness-building does not propagate into AI mediation efficiently. Identity-aligned editorial positioning does, with disproportionate efficiency in categories where consensus is tight. The leverage of the right editorial position in the right CEP is far higher than the leverage of broad-distribution awareness investment.

#### **Pattern 6: AI training data lag creates phantom brand presence**

The sixth pattern emerged from a single category measurement but is sufficiently striking to warrant inclusion as a standalone finding. It describes a structural property of AI-mediated brand visibility that no prior pattern captures: the temporal lag between actual market reality and the AI tier’s representation of that market.

In March 2024, Intuit decommissioned Mint, its consumer personal finance application that had been the dominant brand in the category for over a decade. Existing users were migrated to Credit Karma; the Mint application was retired; the brand effectively ceased to exist as a current consumer product. By April 2026, when this measurement was conducted, Mint had not been a live application for 25 months.

In Third System’s measurement of personal finance applications, Mint was mentioned in 44% of AI responses. Among brands with non-zero Presence scores, Mint placed fifth — ahead of Rocket Money, EveryDollar, Quicken Simplifi, and PocketGuard, all of which are live applications with active user bases. The AI tier mentioned a brand that had not existed for over two years more frequently than it mentioned multiple applications consumers can actually use today.

Per-model breakdown reveals an additional dimension. On Anthropic, Mint surfaced in 65% of responses. On OpenAI, Mint surfaced in 23% of responses. The 42-point spread is the second-largest in the category dataset, and it suggests that the two AI models have meaningfully different training-data freshness on this specific category change. Anthropic’s recommendations remain anchored in pre-shutdown training data more strongly than OpenAI’s. Both systems lag actual market reality, in different ways and to different extents.

The mechanism is straightforward. AI training corpora are necessarily backward-looking. They capture the discourse that existed when the data was crawled, not the discourse that exists at the moment a query is made. For a brand that had been the dominant player for over a decade and had accumulated proportionate volume of editorial coverage, blog posts, comparison articles, and how-to guides, that backward-facing volume continues to surface in responses long after the brand itself has stopped existing. The AI tier does not have a mechanism to remove a brand from its training corpus when the brand shuts down. The corpus simply ages, and the brand persists within it.

The implications follow directly. AI brand visibility is not a current measurement of market reality. It is a measurement of how the AI tier currently weighs an aging corpus against more recent inputs. This makes the AI Presence Index a hybrid signal: partly reflective of current market position, partly reflective of accumulated marketing investment that may extend many years backward, partly reflective of category disruptions that have not yet propagated through the AI's training systems.

For brand strategy, the implications cascade. A brand that has shut down, been acquired, or rebranded continues to generate AI presence for an indeterminate period after the change. Marketers may be allocating budget against a competitor that no longer exists, while their own past marketing investment continues to generate AI presence after their products are discontinued. A brand currently rebranding may face a multi-year period during which the new brand name underperforms in AI mediation while the old name persists. A brand entering a category currently dominated by a defunct brand may find it easier to win AI presence than the leaderboard structure would suggest, because the dominant entity is not actually competing for share.

The single-category basis for this finding limits its current generalizability. Other category disruptions — Twitter's rebrand to X, the GE Brands company splits, the Toys R Us shutdown and limited revival, several pharmaceutical brand transitions — would provide additional test cases. Phase 2 continuation will include at least one category selected specifically for the presence of a recent major brand disruption, to test whether the phantom brand presence pattern observed for Mint replicates in other category-shutdown contexts.

In AI mediation, brands die slowly. The brand value of a discontinued brand persists in the AI tier for years after the brand stops existing in the market. Every brand strategist needs to know whether their category contains such phantoms, because the leaderboard cannot be read accurately without that knowledge.

## Limitations

Five categories meets the threshold at which cross-category claims become substantively defensible. They do not meet the threshold at which the patterns can be claimed as universal. Six limitations apply to the findings reported here.

**Sample size.** Each category measurement comprises 96 successful API calls, producing approximately ten-percentage-point confidence bands around individual brand Presence scores. The cross-category patterns reported here survive that uncertainty because the patterns are large relative to the band. Smaller margins between specific brands within a category should not be treated as decisive.

**Two models, not three.** The original measurement design included Google's Gemini, which has been excluded from v0.3, v0.4, and v0.6 results because of a Workspace-domain access restriction during the study window. Gemini will be reinstated when the restriction clears.

Cross-model variance reported here is likely a lower bound; introducing a third model historically widens the range, and the variance findings would likely strengthen.

**Five categories, not many.** Five categories is enough to claim replication of patterns across diverse contexts. It is not enough to claim that the patterns hold in general. Phase 2 continuation will introduce additional categories specifically designed to test framework predictions, including a category designed to test the discourse-language bias hypothesis introduced in Pattern 4.

**A single point in time per category.** Each category was measured on a single day in late April 2026. The findings describe AI behavior on those days. They do not describe how AI brand visibility moves over time. AI models are updated frequently, often without external notice. Longitudinal measurement requires periodic re-baselining and will be added in Phase 2 continuation. The phantom brand presence pattern (Pattern 6) is the most time-sensitive of the findings; its strength may evolve as AI training data refreshes.

**Construct validity remains unproven.** The AI Presence Index measures a real and stable property of the AI tier, but whether that property correlates with consumer consideration, purchase intent, or sales is currently unknown. A correlation study against external brand-tracking data is planned for Phase 3. Until that study is complete, AI Presence should be treated as a leading indicator with unverified predictive value.

**Registry construction is not perfectly neutral.** The brand registry for each category is curated and reflects the framework's view of which brands constitute the category's competitive set. Two of the five categories required registry revision after the first measurement surfaced significant unknown brands. The unknown-mentions list provides a continuing diagnostic, but registry construction remains the methodological choice with the most direct effect on which brands appear in measurements.

## What's next

Phase 2 will continue with two additional category measurements designed to test specific framework predictions. The first is a category selected explicitly to test the discourse-language bias hypothesis from Pattern 4: a category in which a non-English-discourse country dominates global production but has limited US-targeted English-language marketing. Candidate categories include Japanese kitchen knives, French wine, and Korean small electronics. The second is a category selected to test the phantom brand presence hypothesis from Pattern 6: a category with a recent major brand disruption (shutdown, acquisition, or rebrand) sufficient to generate measurable training-data lag effects. Candidate categories include social media platforms (testing the Twitter-to-X transition), retail (testing the Bed Bath & Beyond shutdown and revival), and several pharmaceutical brand transitions.

Phase 3 is the construct-validity study. The structure is straightforward. AI Presence scores from at least three categories, measured at two points in time, will be correlated against consumer-tracking data from the same windows. The result will determine whether AI-mediated visibility predicts shifts in consideration, purchase intent, or share of category. A finding of strong correlation establishes the AI Presence Index as a leading indicator with predictive value beyond the AI tier itself. A finding of weak correlation establishes the index as a measurement of AI behavior alone, with strategic implications limited to AI-channel optimization. Both outcomes are publishable, and both are useful.

Phase 4, contingent on Phase 3 results, will release the remaining five AIAS components: Ranking, Consistency, Coverage, Grounding, and Sentiment. Each will be added as a sep-

arate measurement layer, with documented rationale and a defined trigger condition for inclusion. The order of release will follow the priority specified in the methodology log: Consistency first because it is computable from existing data, Ranking second because the framework's strategic implications depend on it, the others as cross-category data permits.

Third System publishes its methodology, its data, and its corrections openly. Each subsequent report includes the underlying response dataset alongside the leaderboard, allowing critics, peer reviewers, and customers to verify the analysis. Brand strategy in the AI-mediated era requires measurement that can be inspected.

---

*Authored by Pablo Ulpiano Gonzalez Castro. Methodology log: methodology.md v0.3. Underlying datasets: presence\_index\_v0.3\_pmsoftware.csv, presence\_index\_v0.3\_runningshoes.csv, presence\_index\_v0.3\_oliveoil.csv, presence\_index\_v0.3\_skincare.csv, presence\_index\_v0.3\_finance.csv. Inquiries: hello@thirdsystem.ai.*