

Measuring Presence Twice

A pre-registered test of AI presence consistency, and why it could not be measured as proposed.

June 2026 · Pablo Ulpiano González Castro

AIAS™ Measurement Program · Methodology v1.7 · Third System™

The programme set out to add a second dimension to AI brand presence: not how present a brand is, but how evenly. The dimension collapsed into the first. At the mention volumes today's models produce, the evenness of a brand's AI presence cannot be told apart from its level.

Brand teams face a widening menu of AI-visibility metrics, among them scores that claim to measure the consistency or stability of a brand's presence across models and prompts. This study tested whether such a measure carries information that presence does not. Under a locked, pre-registered protocol, it does not. The reason is arithmetic rather than incidental, and it applies to any consistency score built the same way.

score built from that spread restates the average. A second pattern surfaced along the way: many brands that every model recognizes draw almost no spontaneous mention, a gap widest among prestige labels. No consistency measure is adopted on this evidence. The measure returns for redefinition, and practitioners should read consistency scores of this construction as presence under another name.

Consistency was defined as the evenness of a brand's recall across a fixed six-model panel and tested against a requirement set in advance: to count as a new measure, it had to carry information beyond presence. It did not. Consistency rose with presence in every category tested, at a pooled rank correlation of 0.77, well above the threshold fixed beforehand. The cause is a property of small counts. Where mentions are few, the spread of a count is governed by its average, so a

What we measured

Twenty-four brands in each of three categories (skincare, cosmetics, and automotive) were put to a fixed panel of six language models. Each model was asked in a category framing and a cultural framing, and a brand's recall was the count of mentions it drew. A consistency score was formed from the spread of those counts across the six models, higher where a brand read evenly and lower where it scattered. A rule fixed before any score was computed required the consistency score to move independently of presence; a close correlation would mark it as a restatement rather than a measure. Defunct brands were included as a control and were expected to return no reading.


FINDING 01

Consistency followed presence

CPC tracks Presence — it does not dissociate

Consistency rises monotonically with presence level across all three substrates

At recall counts $r \in 0-6$, $CV \approx 1/\sqrt{\text{mean}}$, so CV-based consistency reparametrizes level.



H_CPC_Dissociates FALSIFIED — $|\rho| = 0.77 \geq 0.50$ ceiling; redefinition escalated to v1.8.

Source: AIAS™ v1.7 (CPC.r3) | Protocol v1.7-prereg-r2 | Third System™

Finding 1 · Consistency rose with presence, not apart from it. Each brand’s consistency score against its presence, across skincare, cosmetics, and automotive. The two move together (pooled rank correlation 0.77), so a consistency score of this construction carries no information that presence does not.

In every category, brands with stronger presence scored as more consistent and brands with weaker presence as less so. Pooled across the three categories the rank correlation was 0.77, and it was positive in each one. The two quantities moved together rather than apart.


FINDING 02

The link is arithmetic, not incidental

Spread is governed by the mean

Coefficient of variation against mean recall, with the Poisson prediction

At the counts models produce, CV approaches $1/\sqrt{\text{mean}}$ — so a CV-based score reparametrizes level.



The presence-consistency coupling is arithmetic, not incidental — it recurs for any low-count CV score.

Source: AIAS™ v1.7 (CPC.r3) | Protocol v1.7-prereg-r2 | Third System™

Finding 2 · The link is arithmetic. Each brand's coefficient of variation against its mean recall, beside the Poisson curve on which spread falls as the square root of the mean. At the counts the panel produces the points track the curve, so the spread a consistency score is built from is set by the average it would otherwise add to.

At the mention counts the panel produced, the spread of a count is governed by its average: rare mentions vary little in absolute terms, frequent ones more. A score built from spread divided by average therefore tracks the average. The coupling


is a feature of counting at low volumes, so it would recur for any brand set and any consistency score of the same construction, including those sold by visibility tools.

FINDING 03

Known is not the same as named

Many in-market brands have no defined CPC

Defined vs N/A under the locked floor (mean combined recall $r \geq 1.0$)
The N/A brands are recognized but rarely recalled (e.g. Lancôme, Chanel, Dior at $C_P=6$, ~ 0 recall).



H_CPC_Defined FALSIFIED — pooled defined-rate 57% (43% of in-market brands N/A).
Source: AIAS™ v1.7 (CPC.r3) | Protocol v1.7-prereg-r2 | Third System™

Finding 3 · Recognized is not recalled. Share of in-market brands with a usable consistency score, by category. Many brands every model recognizes draw too little spontaneous mention to score — the gap widest among prestige labels — so a recall-based reading is undefined for them.


Many brands recognized by every model in the panel drew almost no spontaneous mention when a buyer's question invited one. The gap was widest for prestige and heritage labels; functional and category-defining brands held the recalled set. Recognition proved necessary but not sufficient for presence in AI answers.

FINDING 04

Absent brands read as absent

Phantom brands return undefined CPC, as predicted

Cell-D defunct automotive brands: zero recall → below floor → N/A
Consistency is undefined for brands with no AI presence to be consistent about.



H_CPC_PhantomNull CONFIRMED — all 5 Cell-D defunct brands → N/A.
Source: AIAS™ v1.7 (CPC.r3) | Protocol v1.7-prereg-r2 | Third System™

Finding 4 · Absent brands read as absent. Defunct brands drew no mention across the panel and returned no consistency score, shown beside the in-market recall spread for reference — the boundary behaving as designed. Defunct brands drew no recall and returned no consistency no reading rather than a misleading one. score, as the design intended. A brand with no presence yields

Limitations

The finding is bounded to the three categories for which the panel measure exists, and to the mention volumes the current design produces. A denser set of prompts would raise those volumes and might loosen the coupling, so the result states that the measure fails at present volumes, not that presence stability is unmeasurable in principle. Recognition stood at its ceiling for every scored brand, which prevented a direct test against recognition alone; the conclusion rests instead on the measured coupling between the consistency score and mention level, which is unambiguous.

What's next

Consistency returns for redefinition on a measure that is independent of level at the volumes a model panel produces. Candidate measures exist and will each face the same independence test before any is adopted. The recognition–recall gap is taken up as a line of work in its own right, since it bears on how a brand known to AI becomes a brand named by it. No consistency score is shipped until a measure clears the test.

What this means for practitioners

Five propositions, drawn from the locked result and written for brand teams rather than reviewers.

#	Proposition	Status
P1	At current mention volumes, the consistency of a brand's AI presence cannot be separated from its level.	Established
P2	Consistency or variance scores computed as spread over mention level restate presence; they add no second diagnostic.	Established
P3	Recognition by AI is not recall by AI; a brand every model knows can still go unnamed.	Established
P4	The recognition–recall gap is widest for prestige and heritage brands; functional brands tend to own recall.	Observed
P5	No defensible measure of AI presence stability exists yet; treat any AI-consistency score as presence relabelled until one passes an independence test.	Open

The propositions in full

Across skincare, cosmetics, and automotive, the consistency score and the presence measure moved together at a pooled rank correlation of 0.77. A brand's evenness across models is, in practice, a function of how present it is. There is no separate consistency lever to pull.

The correlation is not a quirk of these brands. Spread over level is mechanically tied to level when counts are small, which they are throughout model output. Any score of this construction, whoever supplies it, carries the same defect: it reports presence a second time.

Recognition and recall came apart in the data. A brand can be identified by every model in the panel and still almost never be offered when a category or cultural question is asked

unprompted. In AI-mediated discovery the second is what reaches the buyer.

The brands most exposed are those whose standing rests on being known. Prestige and heritage labels were recognized everywhere and recalled rarely, while functional and category-defining brands occupied the recalled set. An awareness-led position is the one most at risk when answers, not shelves, do the surfacing.

Stability of AI presence may yet be measurable, but not as specified here. Until a measure demonstrates independence from presence at realistic mention volumes, a consistency score should be read as presence under another label, and bought or built accordingly.

The measure was specified, locked, and tested in the open, and it failed a rule set before the data were seen. That is the method working, not faltering. The result narrows what a consistency measure can be, warns against scores already on sale, and leaves the programme with a sharper question than the one it began with.

About the Third System

The Third System describes a shift in how consumers make decisions: instead of browsing and comparing options directly, many now rely on AI systems to generate recommendations. In this environment, brands compete not only for human attention but for inclusion within AI-generated outputs. This creates a new competitive layer where visibility is determined by how AI systems interpret and surface brands.

Authored by

Pablo Ulpiano González Castro

School of Visual Arts, MPS Branding Program, New York, NY

Third System™ (research entity)

Methodology

This report presents findings from Third System's AI Availability research program. The AI Availability Score (AIAS) is an early-stage measurement framework introduced in Tri-System Brand Growth (Gonzalez Castro, 2026). Findings should be used directionally. Construct validity — the correlation between AI Availability and consumer behavior — has not yet been established and is the subject of Phase 3 of the research program. Methodology is published in full and versioned at thirdsystem.ai/methodology (current: v1.7). Registry construction is curated and reflects the framework's view of each category's competitive set; revisions are documented in the methodology log.

Citation

González Castro, P. U. (2026). *Consistency without Independence: A Pre-Registered Test of Coefficient-of-Variation as the Consistency Component of AIAS (v1.7)*. Third System (SSRN 6878818). thirdsystem.ai/methodology/v1_7

Underlying datasets

Per-brand CPC: osf.io/ec6wh/methodology/v1_7/data/v1_7_cpc.csv

Verdicts: osf.io/ec6wh/methodology/v1_7/v1_7_cpc_verdicts.json

Scoring code: osf.io/ec6wh/methodology/v1_7/scoring/score_v1_7.py

Methodology log: v1.7 (SSRN 6878818).

Inquiries: hello@thirdsystem.ai