

How Evenly, Not Just How Much

A pre-registered redesign of AI presence consistency, and why it now reads on its own, apart from how present a brand is.

June 2026 · Pablo Ulpiano González Castro

AIAS™ Measurement Program · Methodology v1.8 · Third System™

The programme set out to add a second dimension to AI brand presence: not how present a brand is, but how evenly. An earlier attempt collapsed into the first. At the mention volumes today's models produce, the evenness of a brand's presence could not be told apart from its level. The redesigned measure separates them, and evenness now reads on its own.

Brand teams face a widening menu of AI-visibility metrics, among them scores that claim to measure the consistency or stability of a brand's presence across models and prompts. An earlier study found that the obvious way of building such a measure restated presence rather than adding to it. This study redesigns it. Under a locked, pre-registered protocol, the new measure moves independently of presence level, reaches brands the old one could not score, and tells how evenly a brand surfaces apart from where it surfaces. Consistency becomes a second reading, not a second name for the first.

Consistency was defined as the evenness of a brand's presence across a fixed six-model panel and tested against a

requirement set in advance: to count as a measure, it had to carry information that presence does not. An earlier version failed that test. It rose with presence in every category, at a pooled rank correlation of 0.77, because at low mention counts the spread of a count is governed by its average, so a score built from that spread restates the average. The redesigned measure removes the coupling by construction. Instead of dividing a brand's spread by its average, it weighs the spread the models actually show against the spread chance alone would produce, which cancels the level term. On the same brands, its correlation with presence falls to 0.091, far below the line fixed beforehand, against 0.682 for the measure it replaces. It also reaches further: twenty-nine brands recalled too rarely for the old measure to score at all now carry a reading, with no failure at the opposite, saturated end. A companion measure reads where a brand surfaces rather than how much, and moves separately enough to count as a distinct facet. Consistency is now a measurable second component of AI presence. The next step is combining the two readings into one; until then, a brand's evenness can be read as genuine information about its AI presence, no longer a restatement of its level.

What we measured

The redesigned measure was tested not on fresh data but on an archive the programme had already collected and frozen: five categories, headphones, skincare, cosmetics, automotive, and premium spirits, each put to the same panel of six models, one hundred twelve brands in all. Holding the data fixed puts the test on the instrument rather than the brands, since the same records the earlier measure failed on were re-read under the redesign. For each brand, the six models' mention counts were set against what chance alone would produce if every model surfaced the brand at one underlying rate; a

brand whose models agree beyond chance reads as consistent, one whose models disagree as inconsistent. A rule fixed before any value was computed required the new measure to move independently of presence, the test the earlier measure failed. Two companion readings were specified alongside it: one for where a brand surfaces rather than how evenly, and one for graded recognition depth, the second held for a later collection because the archived recognition signal is recorded only as yes or no.


FINDING 01

Evenness now reads on its own

Dispersion that does not restate the mean

ϕ is independent of recall level; the CV-based measure it replaces is not

Same six-model panel, same brands — the contrast is the instrument, not the data.



H_MeanIndependent CONFIRMED — pooled $|\rho(\phi, \mu)| = 0.091$, below the 0.50 ceiling; CV-CPC baseline 0.682.

Source: AIAS™ v1.8 (CPC redesign) | Protocol v1.8-prereg-r4 | Third System™

Left: the redesigned measure plotted against presence level, flat across the range at a correlation of 0.091. Right: the measure it replaces, against the same level, sloping at 0.682. The same brands sit on both sides; what changes is the instrument, not the data.

The redesigned measure no longer tracks how present a brand is. Across the archive its rank correlation with presence level is 0.091, close to none, against 0.682 for the measure it replaces, read on the very same brands. The earlier coupling was not a

fact about brands but about how the measure was built; rebuilt to weigh each brand's spread against what chance alone would produce, it carries no level term, and evenness becomes a reading in its own right.


FINDING 02

It reaches the brands the old measure could not score

 ϕ extends defined coverage into the low-recall floor

The 29 brands CV cannot score — recalled too rarely — still yield a finite ϕ

ϕ is undefined only at true-zero and full saturation; CV is undefined below mean recall 1.0.



H_LowRecallDefined CONFIRMED [0] — ϕ -defined strictly contains CV-defined; non-empty gain; 0 ceiling units.

Source: AIAS™ v1.8 (CPC redesign) | Protocol v1.8-prereg-r4 | Third System™

Where each measure can be computed across the range of recall. The redesigned measure covers all eighty-four brands that surface at all, twenty-nine of them recalled too rarely for the older measure, which reaches only fifty-five. The gain is exactly the sparse, low-recall band a stability reading most needs.

Twenty-nine brands are recalled so rarely that the earlier measure could return nothing for them, and these are exactly the brands a stability reading is most useful for. The redesigned measure scores all of them, and does so without

breaking at the opposite extreme, where a brand is named by every model every time. The reading now spans the full range of presence, sparse to saturated, rather than only the well-recalled middle.


FINDING 03

How evenly and where are different questions

 ϕ and J measure distinct facets of consistency

How much a brand surfaces (ϕ) and where it surfaces (J) come apart

Below the 0.70 redundancy ceiling, with discordant brands present in every framing.



H_PositionalDissociation CONFIRMED — $|\rho(\phi, J)| = 0.403$, below the 0.70 ceiling; 9 discordant.

Source: AIAS™ v1.8 (CPC redesign) | Protocol v1.8-prereg-r4 | Third System™

How evenly a brand surfaces, plotted against where it surfaces, one point per brand. The two readings move apart at a correlation of 0.403, well below the line at which one would be redundant; the nine ringed brands sit steady on one reading and scattered on the other.

A brand can be surfaced at a steady rate across the panel yet in scattered places, or named in the same places by models that disagree on how often. The redesign reads the first; a companion measure reads the second; and the two move

apart enough, a correlation of 0.403, well below the line for redundancy, that nine brands sit consistent on one facet and inconsistent on the other. How much and where are genuinely two readings, not one.


FINDING 04

The reading is stable across measurements

ϕ rank order is moderately stable across waves

Per-brand ϕ at two acquisition waves of the same panel

Reported as moderate ($\rho = 0.65$), not overstated — a single re-acquisition, not a stability claim.



Two-wave reproducibility: $\rho(\phi) = 0.645$ across waves t_1 to t_2 ($n = 83$); supplementary, non-gating.
Source: AIAS™ v1.8 (CPC redesign) | Protocol v1.8-prereg-r4 | Third System™

Each brand's evenness at a first measurement against a second, on the same panel. The points cluster along the line of equality at a correlation of 0.645, moderate, enough to show the ordering is not an accident of a single draw.

Re-measured at a second wave, the brands' order on the new measure holds at a rank correlation of 0.645, moderate, and reported as such rather than overstated. The ordering is not an artefact of a single draw, though one re-measurement is not yet the evidence a high-stakes index would need. Enough to trust the reading, not yet enough to stop checking it.

Limitations

The test rests on an archive of five categories, re-read rather than freshly collected, and read in three overlapping views rather than three independent ones; two of the views coincide on the same brands, so the headline figure summarizes the evidence rather than replicating it. The counts behind the secondary findings are modest: the coverage gain rests on twenty-nine brands, the split between how much and where on nine. The archived surfacing was recorded by three

different methods across the categories, uniform within each but not across them, which the design accounts for and a single-method collection would tighten. Graded recognition depth could not be scored, because the archived recognition signal is yes or no only, and is held for a later collection. And the measure validated here is one component of a larger score, not the whole; its soundness does not yet establish the soundness of any combination.

What's next

With a measure of how evenly that holds and a measure of where beside it, the open task is composition: combining the two into a single consistency reading without letting the level dependence back in, and settling how they trade off where they disagree. That combination is the next pre-registered step. Graded recognition depth rejoins the programme at the first collection that records more than yes or no. Beyond

consistency, the remaining parts of the score, ranking, coverage, grounding, and sentiment, each face the same build-and-validate arc this phase completed here before any of them enters the composite. The structure work the earlier, broken measure could not support is reopened now that a sound one exists.

What this means for practitioners

Five propositions, drawn from the locked result and written for brand teams rather than reviewers.

#	Proposition	Status
P1	The evenness of a brand's AI presence can now be measured on its own, apart from how present the brand is.	Established
P2	The redesigned measure reads the sparsely recalled brands the earlier one could not score, where a stability reading matters most.	Established
P3	How evenly a brand surfaces and where it surfaces are distinct readings; a brand steady on one can be scattered on the other.	Established
P4	The reading holds across re-measurement: stable enough to track a brand over time, not yet stable enough to stand as a high-stakes single number.	Observed
P5	Consistency is a validated component, but the score that combines it with the others is not yet built; read one consistency figure as an input, not a finished verdict.	Open

The propositions in full

Read on the same brands the earlier measure restated presence on, the redesigned measure's correlation with presence level is 0.091, against 0.682 before. The difference is the construction, not the data: weighing each brand's spread against what chance would produce cancels the level term that dividing by the average had carried. A brand's evenness is now a separate lever from its presence.

Twenty-nine brands surface too rarely for a spread-over-average measure to return anything, and rarely surfaced brands are precisely where a stability reading earns its keep. The redesign scores every one of them, and holds at the other extreme too, where a brand is named by every model every time. The reading spans sparse to saturated rather than only the comfortable middle.

Two brands with the same evenness across the panel can occupy the same handful of prompts or scatter across many. The redesign reads how evenly; a companion reads where. They move together loosely enough, a correlation of 0.403,

below the line at which one would make the other redundant, that nine brands land consistent on one and inconsistent on the other. Carrying both keeps a real distinction that a single number would erase.

Re-measured at a second wave, the brands' order on the new measure recurs at a rank correlation of 0.645. That is moderate, and stated as moderate: enough to say the ordering is not an accident of one measurement, short of the reliability a published single-number index would demand. The reading is trustworthy to track, and still worth re-checking.

The redesign makes consistency a sound input, not a finished score. Combining it with the other parts of AI presence, and with its own positional companion, is a separate problem with its own test, still ahead. Until that composite is built and validated, a consistency figure should be read as one well-made reading among several, not as a brand's standing on its own.

The redesigned measure was specified, locked, and tested in the open against the same rule the earlier one failed, and this time it held. That is the method compounding: a failure that named exactly what a sound measure had to do, followed by a measure built to do it. Consistency is now a second reading of AI presence rather than a second name for the first, and the programme moves from repairing an instrument to assembling a score from instruments that work.

About the Third System

The Third System describes a shift in how consumers make decisions: instead of browsing and comparing options directly, many now rely on AI systems to generate recommendations. In this environment, brands compete not only for human attention but for inclusion within AI-generated outputs. This creates a new competitive layer where visibility is determined by how AI systems interpret and surface brands.

Authored by

Pablo Ulpiano González Castro

School of Visual Arts, MPS Branding Program, New York, NY

Third System™ (research entity)

Methodology

This report presents findings from Third System's AI Availability research program. The AI Availability Score (AIAS) is an early-stage measurement framework introduced in Tri-System Brand Growth (Gonzalez Castro, 2026). Findings should be used directionally. Construct validity — the correlation between AI Availability and consumer behavior — has not yet been established and is the subject of Phase 3 of the research program. Methodology is published in full and versioned at thirdsystem.ai/methodology (current: v1.8). Registry construction is curated and reflects the framework's view of each category's competitive set; revisions are documented in the methodology log.

Citation

González Castro, P. U. (2026). *A Mean-Independent Consistency Instrument for Cross-Model AI Presence: Redesigning CPC on a Quasi-Binomial Dispersion Basis*. Third System (SSRN pending). thirdsystem.ai/methodology/v1_8

Underlying datasets

Per-brand instrument table: osf.io/ec6wh/methodology/v1_8/v1_8_instrument_table.csv

Verdicts: osf.io/ec6wh/methodology/v1_8/v1_8_verdicts.json

Scoring code: osf.io/ec6wh/methodology/v1_8/scoring/score_v1_8.py

Materialization manifest: osf.io/ec6wh/methodology/v1_8/v1_8_materialization_manifest.json

Methodology log: v1.8 - results-locked v1.8-results-r1 (commit 85d2202); SSRN pending.

Inquiries: hello@thirdsystem.ai