Third System
AI Presence Index

The Recognition Ceiling on the Premium Spirits Substrate

Premium spirits is the first AIAS substrate where every brand in the 24-brand registry achieves rich recognition (R3) across all six LLMs. AI knows these brands perfectly. The question that remains is which brands AI recommends — and there, the field splits wide open.

May 2026 · Pablo Ulpiano González Castro · Third System

Independent measurement for the AI mediation layer.

AIAS v0.23 measures AI Presence for 24 premium spirits brands across six large language models under Protocol v1.6. The substrate introduces the seventh product family to the AIAS measurement program and the first where editorial-authority and cultural-cult recall channels were expected to populate with comparable density.

Four findings shape the v0.23 phase. First, all 24 brands scored R3 (rich recognition) across all six models — a universal ceiling never observed in the prior six substrates. Second, recall, not recognition, drives the entire Presence distribution: composite scores range from 40.0 (zero recall) to 71.2, with six brands perfectly known yet never recommended. Third, Fortaleza, an independent craft tequila brand, lands in the Dominant regime alongside Hennessy, Patrón, and Grey Goose — conglomerate ownership confers no statistically significant AI Presence advantage (p = 0.28). Fourth, the editorial-authority and cultural-cult recall channels converge more than in any prior substrate (overlap = 0.80), suggesting that in high-familiarity categories the two channels surface largely the same brand set.

The recognition ceiling is the defining structural feature of this substrate. With all 144 Phase A probes returning R3, the Recognition axis contributes zero discriminating variance to the Presence composite. This directly falsifies H_ILDirect (r = NaN; zero variance precludes correlation) and shifts analytical attention entirely to the Recall axis.

On recall, the 24 brands distribute across all four regimes in a perfectly balanced 6/6/6/6 split. Six brands — Bombay Sapphire, The Balvenie, Casamigos, Compass Box, and Fernet-Branca, plus Jameson at the margin — register zero or near-zero recall despite flawless recognition. These are brands that AI systems can describe in expert detail but never surface in recommendation contexts. The recognition–recall gap is the actionable finding for brand strategists operating in this category.

The substrate-specific hypothesis H_ConglomeratePortfolio is falsified: conglomerate-owned brands (n = 14) average 60.75

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

on the Presence composite versus 58.05 for independents (n = 10), a 4.6% lift that does not reach statistical significance (t = 0.59, p = 0.28; Mann-Whitney U = 72.0, p = 0.46). Training-corpus density from conglomerate media coverage does not translate mechanically into higher AI Presence.

What we measured

Phase A (Recognition): 24 brands × 6 LLMs = 144 single-shot probes at temperature 0.0. Each response scored R0–R3 by an LLM judge (Claude Sonnet 4.5). Substrate Recognition pre-screen (Protocol v1.6) applied; zero brands excluded.

Phase B (Two-channel Recall): 6 category-level probes (3 editorial-authority, 3 cultural-cult) × 6 LLMs = 36 queries. Each response scored for brand mention (binary), slot position (ordinal rank of first mention), and elaboration (LLM judge). 110 total brand mentions detected across 36 responses.

Composite Presence score: 0.4 × Recognition (normalized) + 0.4 × Recall rate + 0.2 × inverse mean slot position. Regime classification by quartile thresholds: Q25 = 58.28, Q50 = 64.22, Q75 = 67.50.

Model panel (fixed from v0.17): Claude Opus 4.5, Claude Sonnet 4.5, GPT-4o, GPT-4o-mini, Gemini 2.5 Flash, Gemini 2.5 Flash Lite.

Registry: 24 brands spanning 12 spirit types, structured in three tiers (Global-dominant, Premium-enthusiast, Craft-cult-emerging) with conglomerate ownership coded binary (top-5 global spirits groups = 1; all others = 0). 14 conglomerate, 10 independent.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

FINDING 01

The recognition ceiling.

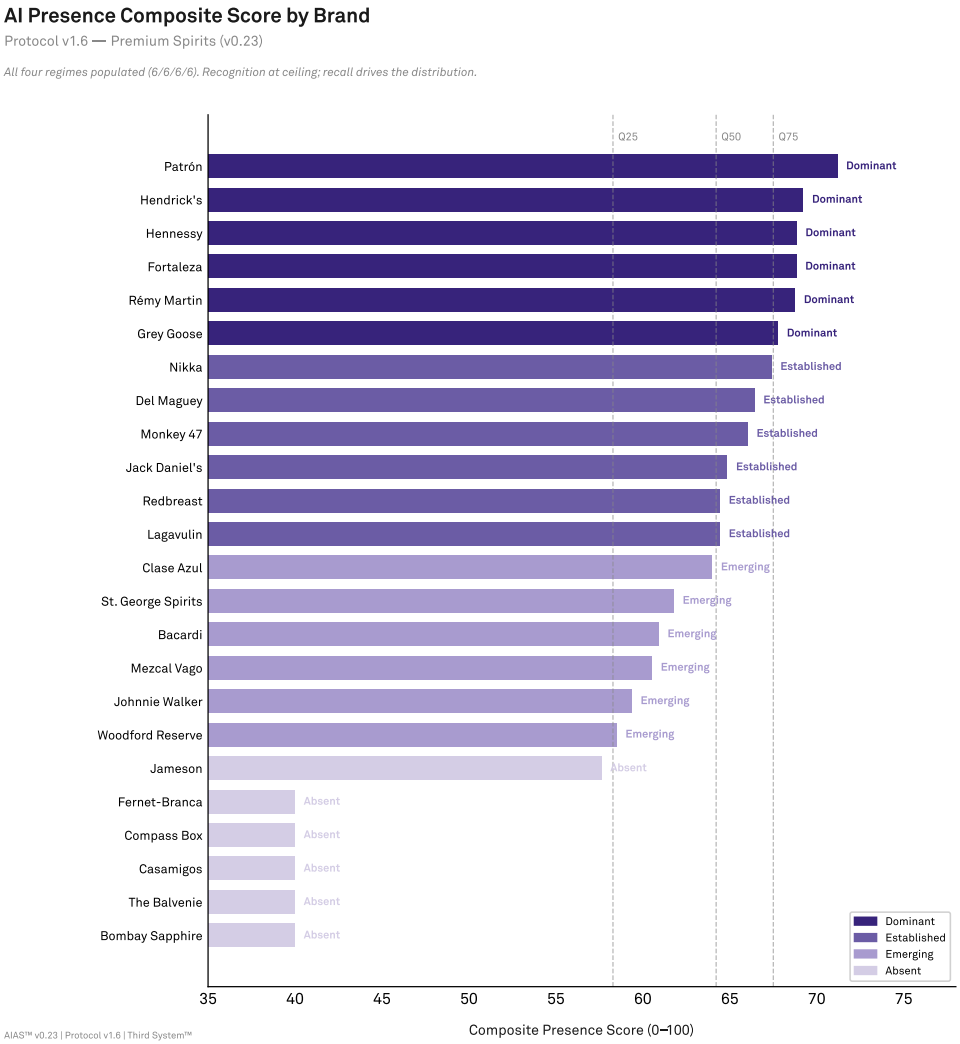


Figure 1 · AI Presence composite score by brand, Protocol v1.6, Premium Spirits (v0.23). Bars are colored by regime; dashed lines mark the quartile thresholds (Q25 = 58.28, Q50 = 64.22, Q75 = 67.50). All four regimes are populated (6/6/6/6). Recognition is at ceiling; recall drives the distribution.

All 24 brands scored R3 across all six models — 144/144 probes at the maximum recognition level. This is structurally distinct from every prior AIAS substrate, where recognition variance contributed meaningfully to the composite. In premium spirits, LLMs can identify distillery locations, flagship expressions, ownership lineage, and tasting profiles for every brand in the registry, from Johnnie Walker to Mezcal Vago. The result: the entire Presence distribution is driven by who gets recommended, not who gets recognized.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

FINDING 02

Known but never recommended: the recognition–recall gap.

Six brands score a composite of 40.0, the floor value representing perfect recognition with zero recall: Bombay Sapphire, The Balvenie, Casamigos, Compass Box, and Fernet-Branca (plus Jameson at 57.63 with a single mention).

These brands exist fully in AI's knowledge base but are invisible in its recommendation behavior. For brand managers, this gap is the most directly actionable finding: recognition alone is a necessary but insufficient condition for AI Presence.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

FINDING 03

Craft disrupts the hierarchy.

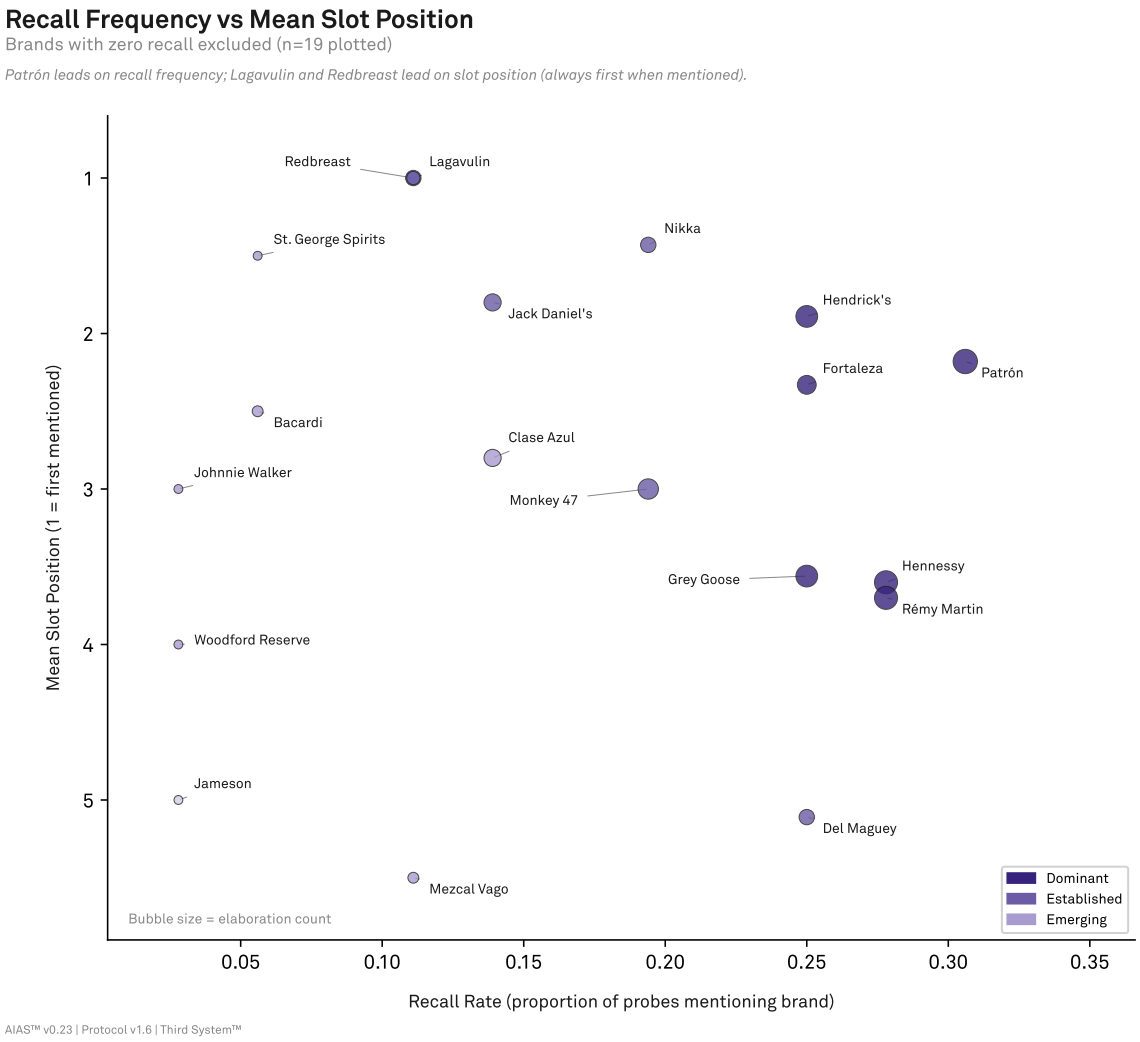


Figure 2 · Recall frequency vs mean slot position (1 = first mentioned), colored by regime. Brands with zero recall are excluded (n = 19 plotted); bubble size indicates elaboration count. Patrón leads on recall frequency; Lagavulin and Redbreast lead on slot position (always first when mentioned).

The Dominant regime (composite ≥67.50) contains two independent brands: Fortaleza (68.84, craft tequila) and Hendrick's (69.23). Fortaleza outranks every Bacardi Ltd. brand in the registry despite being a fraction of the size. Rémy Martin (68.76), coded independent under the top-5 rule, also lands Dominant. The finding is consistent with H_ConglomeratePortfolio's falsification: scale and media spend do not mechanically translate into AI recommendation behavior.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

FINDING 04

Channel convergence at high familiarity.



Figure 3 · Recall mentions by channel and brand: editorial-authority (EA) vs cultural-cult (CC), sorted by total recall. Overlap = 0.80. The channels converge in this high-familiarity substrate; channel-exclusive brands (red outlines), such as Rémy Martin (10 EA / 0 CC) and Jack Daniel's (0 EA / 5 CC), are exceptions.

The editorial-authority channel surfaced 15 brands; cultural-cult surfaced 16; 12 appear in both (overlap = 0.80). This is the highest overlap observed across seven substrates and sits at the exact boundary of H_C2's support range. The finding suggests that in categories where LLMs have dense training data, the two channels converge: the same brands surface whether the query frames authority through critics or through community enthusiasm. Channel-specific brands do exist — Rémy Martin is pure editorial (10 EA / 0 CC), Jack Daniel's pure cultural (0 EA / 5 CC) — but they are exceptions.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

FINDING 05

Conglomerate ownership is not a Presence driver.

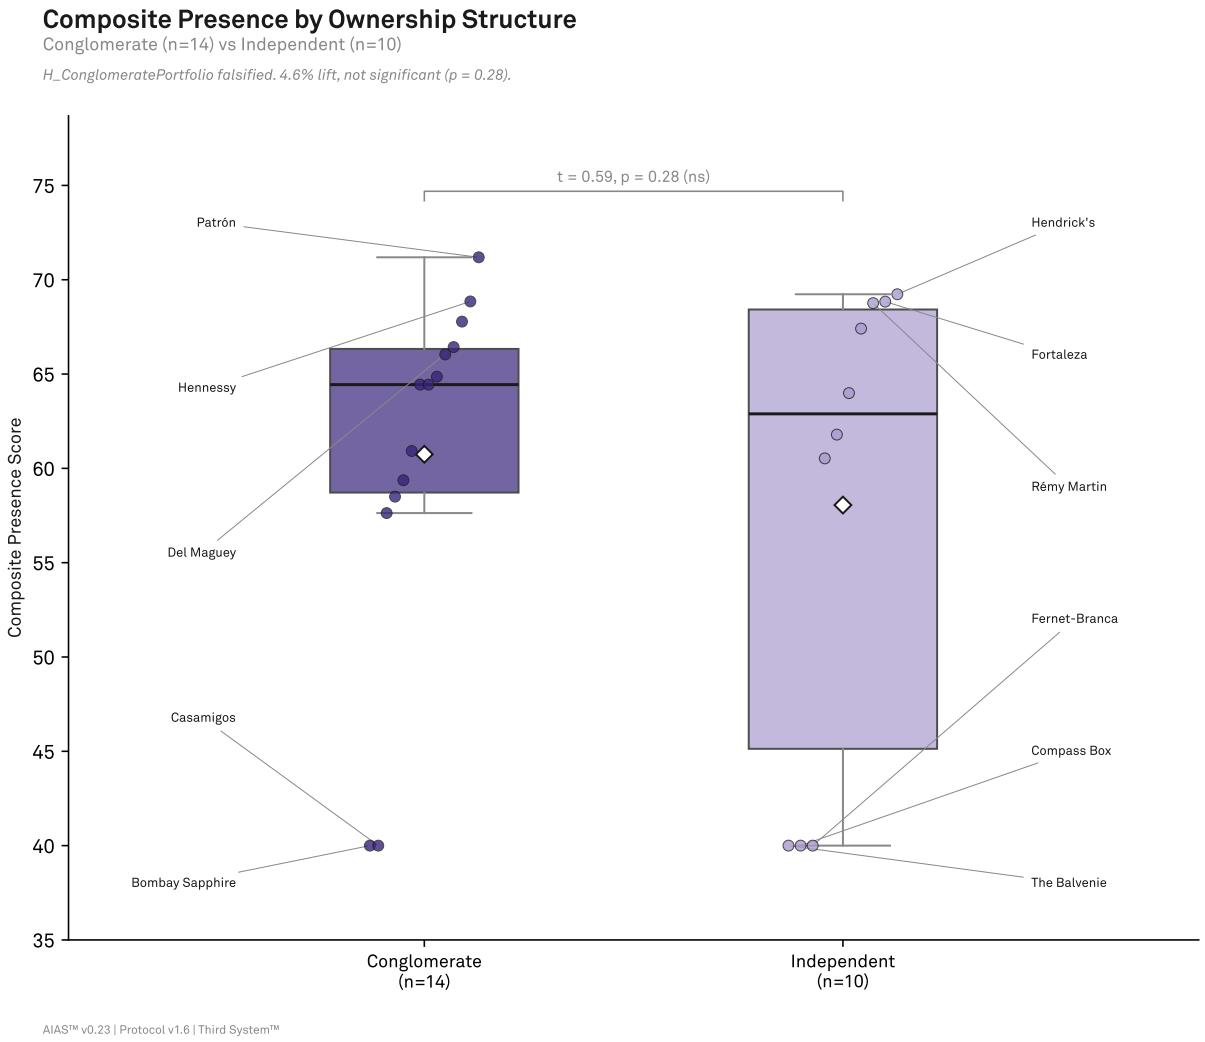


Figure 4 · Composite Presence by ownership structure: conglomerate (n = 14) vs independent (n = 10), brand points overlaid. H_ConglomeratePortfolio falsified: a 4.6% lift that is not significant (t = 0.59, p = 0.28).

The 14 conglomerate-owned brands average 60.75 on the composite versus 58.05 for the 10 independents — a directionally positive but non-significant 4.6% lift (t = 0.59, p = 0.28). Diageo, Pernod Ricard, Bacardi Ltd., LVMH, and Brown-Forman collectively own 14 of 24 registry brands but capture no systematic AI Presence advantage. The mechanism hypothesized — training-corpus density from conglomerate-level media coverage — either does not operate as theorized or is offset by the editorial and enthusiast attention that craft and independent brands attract.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

Limitations

The universal R3 ceiling collapses the Recognition axis and renders H_ILDirect structurally untestable (zero-variance correlation returns NaN). This is a property of the substrate, not a protocol deficiency, but it limits the composite's discriminating power to a single component.

H_C2_TwoChannel's overlap coefficient lands at exactly 0.80, the upper boundary of the pre-registered support range. A stricter threshold (e.g., 0.75) would flip the verdict. The result should be interpreted as borderline rather than robust.

Conglomerate coding uses a binary top-5 rule. William Grant & Sons, Rémy Cointreau, and Asahi Group are coded independent despite being multinational producers. Alternative coding schemes (e.g., any publicly traded parent) could shift H_ConglomeratePortfolio's direction or significance.

The 24-brand registry is a purposive sample. Pan-spirits breadth comes at the cost of within-type depth: some spirit categories have only 1–2 representatives, limiting subcategory-level inference.

Recall probes name specific publications and competitions (Whisky Advocate, SFWSC). This anchoring may favor brands prominent in those specific circuits over brands prominent in others.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

What’s next

The recognition ceiling finding motivates a cross-substrate comparison: which product families exhibit ceiling effects and which show recognition variance? This pattern becomes testable at scale as the substrate count grows.

The recognition–recall gap invites intervention-design work. If a brand like The Balvenie is perfectly known but never recommended, what content or signal changes would shift its recall behavior? This is the bridge from measurement to strategy.

Channel convergence at 0.80 overlap raises the question of whether a third channel — perhaps a commerce or availability-framed prompt — would introduce additional discriminating variance in high-familiarity substrates.

Premium spirits reveals what happens when AI recognition saturates: the competitive frontier shifts entirely to recall. A brand that AI can describe in expert detail but never recommends occupies a new kind of strategic gap — one invisible to traditional brand-tracking methods that measure awareness without distinguishing recognition from recommendation. The AIAS framework surfaces this gap by design.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

Pre-registered hypothesis scoring

Six hypotheses were pre-registered at tag v0.23-prereg-r1, before any acquisition. Falsification criteria were locked in the pre-registration. The verdicts below are what the data produced against those criteria.

H	Pre-registered prediction	Result	Status
H_Regime4	**Standard set.** Each of the 24 brands classifies into one of four Presence regimes (Dominant, Established, Emerging, Absent) on the composite score. Falsified if any regime is empty.	6/6/6/6 split across Dominant, Established, Emerging, Absent.	**SUPPORTED**
H_RecognitionPrescreen	**Standard set.** The R0 pre-screen gate excludes at most three brands from the composite denominator. Falsified if more than three are excluded.	0 brands excluded (threshold: ≤3).	**SUPPORTED**
H_Phantom	**Standard set.** No phantom brands surface. The registry holds no defunct brands, so the hypothesis is carried for protocol completeness, not tested.	No defunct brands in registry.	*NOT TESTABLE*
H_ILDirect	**Standard set.** Information Landscape Direct scores (from Phase A Recognition richness) correlate positively with Phase B Recall frequency ($r \geq 0.40$, $p < .05$). Falsified if $r < 0.40$ or $p \geq .05$.	r = NaN (zero recognition variance; all R3 ceiling).	FALSIFIED
H_C2_TwoChannel	**Standard set.** The editorial-authority and cultural-cult channels surface partly overlapping, non-identical brand sets: an overlap coefficient between 0.30 and 0.80. Falsified below 0.30 or above 0.80.	Overlap = 0.80 (EA: 15 brands, CC: 16, intersection: 12).	**SUPPORTED (BORDERLINE)**
H_ConglomeratePortfolio	**Substrate-specific.** Brands owned by the top-5 global spirits groups (n = 14) score higher on the composite than independents (n = 10), by t-test or Mann-Whitney U at the .05 level. Falsified if $p \geq .05$ or the direction reverses.	Conglomerate mean 60.75 vs independent 58.05; lift 4.6%, p = 0.28.	FALSIFIED

Hypothesis interpretation

Each verdict carries interpretation beyond the pass/fail routing. The verdicts are what the data produced; what each one means for the construct is what the prose below addresses.

H_Regime4. The four-regime taxonomy distributes evenly across the 24-brand registry: Dominant (Hennessy, Patrón, Grey Goose, Hendrick's, Rémy Martin, Fortaleza), Established (Jack Daniel's, Lagavulin, Monkey 47, Nikka, Del Maguey, Redbreast), Emerging (Johnnie Walker, Bacardi, Clase Azul, Woodford Reserve, St. George Spirits, Mezcal Vago), Absent (Bombay Sapphire, Jameson, The Balvenie, Casamigos, Compass Box, Fernet-Branca). The 6/6/6/6 split is the most balanced distribution observed across seven substrates. Note that 'Absent' in this context means absent from recall, not from knowledge — every Absent-regime brand scored R3 on recognition.

H_RecognitionPrescreen. Zero brands triggered the pre-screen exclusion gate (R0 across ≥5 of 6 models). This confirms the prediction that premium spirits has high baseline LLM familiarity. The substrate is the cleanest pass through the

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

pre-screen observed to date.

H_ILDirect. Falsified by structural ceiling. All 24 brands scored recognition mean = 3.0 (R3 across all six models). With zero variance on the Recognition axis, the Pearson correlation between Recognition and Recall is undefined (NaN). This is not a measurement failure but a substantive finding: in categories where AI training data is sufficiently dense, recognition saturates and ceases to discriminate. The implication for AIAS methodology is that high-familiarity substrates may require recall-only composite scoring.

H_C2_TwoChannel. Supported at the boundary. The editorial-authority channel surfaced 15 unique brands; cultural-cult surfaced 16; 12 brands appear in both, yielding an overlap coefficient of 0.80 (the upper limit of the pre-registered [0.30, 0.80] support range). Channel-exclusive brands exist: Rémy Martin (10 EA / 0 CC), Del Maguey (8 EA / 1 CC) skew editorial; Jack Daniel's (0 EA / 5 CC), St. George Spirits (0 EA / 2 CC) skew cultural. But the dominant pattern is convergence — in high-familiarity substrates, both channels largely surface the same brand set.

H_ConglomeratePortfolio. Falsified. The 14 conglomerate-owned brands (Diageo 3, Bacardi Ltd. 4, Pernod Ricard 4, Brown-Forman 2, LVMH 1) averaged 60.75 on the Presence composite versus 58.05 for the 10 independent brands. The 4.6% lift is directionally consistent with the prediction but far from significant ($t = 0.59$, $p = 0.28$; Mann-Whitney $U = 72.0$, $p = 0.46$). The independent brands Fortaleza (68.84) and Hendrick's (69.23) both outperform the conglomerate mean. Conglomerate media density — press releases, annual reports, cross-brand coverage — does not translate mechanically into AI recommendation behavior.

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.

About the Third System

The Third System describes a shift in how consumers make decisions: instead of browsing and comparing options directly, many now rely on AI systems to generate recommendations. In this environment, brands compete not only for human attention but for inclusion within AI-generated outputs. This creates a new competitive layer where visibility is determined by how AI systems interpret and surface brands.

Authored by

Pablo Ulpiano González Castro

School of Visual Arts, MPS Branding Program · New York, NY

Third System™ (research entity; data archive and methodology venue)

Correspondence: pablou@pablou.com · ORCID: 0009-0003-8968-9990

Methodology

This report presents findings from Third System's AI Availability research program. The AI Availability Score (AIAS) is an early-stage measurement framework introduced in Tri-System Brand Growth (Gonzalez Castro, 2026). Findings should be used directionally. Construct validity — the correlation between AI Availability and consumer behavior — has not yet been established and is the subject of Phase 3 of the research program. Methodology is published in full and versioned at thirdsystem.ai/methodology (current: v1.6). Registry construction is curated and reflects the framework's view of each category's competitive set; revisions are documented in the methodology log.

Citation

González Castro, P. U. (2026). ***AIAS Presence Index v0.23: Premium Spirits. Recognition Ceiling, Recall-Driven Regime Classification, and the Conglomerate Ownership Hypothesis Across 24 Brands and Six Large Language Models***. Working paper. Third System. Data and pre-registration: osf.io/ec6wh/v23/

Underlying datasets

v0.23 Phase A acquisition (Recognition, n = 144 probes), raw and scored: osf.io/ec6wh/v23/data/v23_phase_a_raw.json, v23_phase_a_scored.json

v0.23 Phase B acquisition (two-channel Recall, n = 36 queries), raw and scored: osf.io/ec6wh/v23/data/v23_phase_b_raw.json, v23_phase_b_scored.json

v0.23 scoring verdicts (six hypotheses, 24 brand records): osf.io/ec6wh/v23/v23_verdicts.json

v0.23 pre-registration (registry, hypotheses, falsification criteria): osf.io/ec6wh/v23/prereg/

Methodology log: v0.23 pre-registration locked at git tag v0.23-prereg-r1 (commit 6b6b047) before acquisition began. Runs under Protocol v1.6 (SSRN 6816340), with v1.5 two-channel Recall (SSRN 6810758) and the v1.2 four-regime taxonomy (SSRN 6761698). Scoring locked at git tag v0.23-scoring-locked (commit 28879cd). Academic companion: working paper, May 2026 (SSRN posting to be confirmed).

Inquiries: hello@thirdsystem.ai

© 2026 Third System. Methodology and underlying datasets: thirdsystem.ai/methodology. Inquiries: hello@thirdsystem.ai. Findings may be cited with attribution.