

Consistency, Category by Category

A cross-category baseline for AI Availability's Consistency reading (CPC)

How stably AI recalls a brand can be measured — and where market dominance stops predicting it.

June 2026 · Pablo Ulpiano González Castro

Third System™ · AIAS™ Measurement Program · v0.31

Generative systems now sit between buyers and brands. The AI Availability program measures where a brand stands inside those systems. This report establishes the first cross-category reading of the program's second component, Consistency: how stably a brand is recalled across AI models, measured across five categories on a single instrument fixed in advance.

A brand can be present in AI answers without being present reliably. One model names it, another does not; one phrasing surfaces it, the next omits it. Consistency captures that reliability as a single reading on a zero-to-one scale, where a brand recalled uniformly across the model panel approaches one and a brand recalled erratically approaches zero. Before the reading can guide anything, two questions have to be answered: does it actually separate brands, and does it behave differently across categories? Across five anchored categories the answer to both is yes, and the categories order themselves in a way that says more about how a category is discussed than about who leads it.

Five propositions carry this baseline. Consistency is measurable: the reading separates brands within every category examined rather than flattening to a single value (P1). It is category-specific: spirits is the least consistent category and skincare the most, with categories shaped by editorial and cultural discourse scoring lower than utilitarian ones (P2). Recognition is not recall: several of the most dominant brands in spirits, Johnnie Walker among them, are known to the models yet barely recalled in the channels that matter, and so carry no consistency reading at all (P3). The reading is reproducible: it reproduces the program's prior instrument exactly where the two overlap, so categories are comparable rather than re-defined study by study (P4). And it is a floor rather than a forecast: the reading is taken across models at a single point in time, and it covers the five categories where the measurement is uniform (P5).

What we measured

Consistency (CPC) is the stability of a brand's recall across a fixed panel of six AI models. For each brand the analysis counts how often each model recalls it across a set of category prompts, then scores the evenness of those counts: uniform recall scores near one, lopsided recall near zero. The reading uses recall only, not recognition, and was computed by rescaling recall data already collected in earlier phases. No

new model queries were issued. A floor handles near-absent brands: any brand a model panel almost never recalls is marked undefined rather than scored as zero, since the stability of an almost-empty signal is not meaningful. Five categories met a common measurement standard and form the comparable set: audiophile headphones, skincare, cosmetics, automotive, and premium spirits.


FINDING 01

Consistency is measurable.

Consistency is well-defined and varies within every category

Brand-level CPC for defined brands, by substrate, with medians (bar) and N

Non-degenerate spread in all five categories — the reading separates brands rather than flattening.



P1 · Established — consistency is well-defined and separates brands within every category.

Source: AIAS™ v0.31 (CPC baseline) | Protocol v0.31-prereg-r1 | Third System™

P1 · Consistency is measurable. Brand-level consistency among defined brands in each category, with category medians. The reading spreads brands across a real range rather than flattening to one value — it discriminates within every category.

Within every category, the consistency reading spreads brands across a real range rather than clustering at one value.

Skincare runs from 0.52 to 0.94 and automotive from 0.64 to

1.00; even the tightest category, headphones, spans 0.57 to 0.69. The instrument distinguishes brands, which is the first thing any measure has to do.


FINDING 02

Consistency is category-specific.

Consistency differs systematically across categories

Distribution of defined-brand CPC by substrate

Spirits least consistent, skincare most — a significant cross-category difference.



P2 · Established — consistency differs systematically across categories; spirits lowest, skincare highest.

Source: AIAS™ v0.31 (CPC baseline) | Protocol v0.31-prereg-r1 | Third System™

P2 · Consistency is category-specific. The distribution of consistency by category, ordered from spirits at the low end to skincare at the high end. The difference is statistically reliable (Kruskal-Wallis $H = 21.51$, $p < 0.001$); categories driven by editorial and cultural discourse score lower than utilitarian ones.

Median consistency differs across the five categories from 0.60 in spirits to 0.80 in skincare, and the difference is statistically reliable (Kruskal-Wallis $H = 21.51$, $p < 0.001$). The


ordering tracks how a category is talked about: categories driven by editorial authority and cultural enthusiasm are recalled less consistently than utilitarian categories.

FINDING 03

Recognition is not recall.

The floor sends near-zero-recall brands to undefined

Defined vs undefined CPC by substrate (omnibus | supplementary)
Undefined are recognized-but-unrecalled and defunct brands — excluded from the distribution, counted here.



P3 · Documented — brands the panel almost never recalls, including category leaders, carry no reading.
Source: AIAS™ v0.31 (CPC baseline) | Protocol v0.31-prereg-r1 | Third System™

P3 · Recognition is not recall. Defined versus undefined brands by category. Brands a model panel almost never recalls fall below the floor and carry no reading — in spirits these include volume leaders such as Johnnie Walker, which dominance alone would not predict.

Consistency is built on recall, and recall does not follow sales. In spirits, several volume leaders are recalled so rarely by the panel that they have no defined reading: Johnnie Walker sits at a mean recall of 0.17, alongside Jack Daniel's, Bacardi, and Jameson. The same recognition–recall gap appeared earlier in prestige skincare. For a brand owner this is the actionable signal: a brand can be universally recognized and still be missing from the answers buyers actually receive.


FINDING 04

The reading is reproducible.

The generalization reproduces v1.7 exactly

Channel-agnostic CPC against v1.7's two-channel CPC, on v0.20/21/22

Every brand lands on $y = x$ — the channel-agnostic unit contains the locked two-channel unit by construction.



P4 · Confirmed — the consistency reading reproduces the program's prior instrument exactly across all shared brands.

Source: AIAS™ v0.31 (CPC baseline) | Protocol v0.31-prereg-r1 | Third System™

P4 · The reading is reproducible. The generalized reading against the prior locked instrument on the categories they share. All seventy-two brands fall on the identity line, so the cross-category numbers are comparable rather than redefined for this study.

The reading is anchored to the program's prior instrument. Recomputed on the categories the two share, it reproduces the earlier values exactly across all 72 brands. The generalization that lets the reading extend to categories with

different channel structures contains the prior definition as a special case, so the cross-category numbers are comparable rather than a fresh definition introduced for this study.

FINDING 05

This is a floor, not a forecast.

The baseline reads consistency across models at one moment, for the five categories where the measurement is uniform. It describes where brands stand now, not whether they will hold,

and it describes recalled brands rather than full category rosters. Those are the boundaries the later components and the longitudinal work are built to extend.

Three patterns

Three patterns stand out. First, consistency is real and uneven within every category; the reading discriminates among brands rather than assigning everyone the same score. Second, categories differ systematically. Ordered by median, consistency runs from premium spirits at 0.60, through headphones at 0.67, cosmetics at 0.74, and automotive at 0.76, to skincare at 0.80. The categories organized around critics and enthusiasts are the least consistent; the utilitarian categories are the most. Third, and most consequential for

brand owners, recall does not follow market share. In spirits the global volume leaders are largely absent from the recalled set: Johnnie Walker, the leading Scotch by volume, is recalled so rarely that it has no consistency reading, and Jack Daniel's, Bacardi, and Jameson sit in the same position. The brands that do carry a reading mix large premium houses with smaller critical favorites. Being known and being named are different things, and consistency measures the second.

Limitations

This is a baseline, and its claims are bounded. The reading is taken across models within a panel, not across the different AI platforms a buyer actually uses, and it reflects a single point in time rather than a trend. Because consistency is defined only for brands with enough recall to measure, each category's

reading describes its recalled brands, not its full roster. And the comparable set is five categories; two more were observed on a coarser grid and two were measured off-standard, so they are not part of the cross-category comparison.

What's next

Three extensions follow. A second measurement wave would turn this static reading into a test of whether consistency holds over time and across model version changes. A cross-platform reading would separate stability inside one model family from stability across the systems buyers

encounter. And the recognition–recall gap, now seen in spirits and earlier in prestige skincare, is worth treating as a diagnostic in its own right: a way to find brands that are widely known yet quietly missing from the answers buyers receive.

The five propositions, scored

Each proposition is the managerial reading of a locked result, scored for brand teams rather than reviewers.

#	Proposition	Status
P1	Consistency is measurable.	Established
P2	Consistency is category-specific.	Established
P3	Recognition is not recall.	Documented
P4	The reading is reproducible.	Confirmed
P5	This is a floor, not a forecast.	<i>By design</i>

Consistency joins Presence as a measurable coordinate of a brand's standing in AI. The cross-category baseline shows the reading works, that it varies with the character of a category, and that it surfaces a gap dominance alone hides: being known is not being named. That gap is where the next questions, and the next opportunities for brand owners, begin.

About the Third System

The Third System describes a shift in how consumers make decisions: instead of browsing and comparing options directly, many now rely on AI systems to generate recommendations. In this environment, brands compete not only for human attention but for inclusion within AI-generated outputs. This creates a new competitive layer where visibility is determined by how AI systems interpret and surface brands.

Authored by

Pablo Ulpiano González Castro

School of Visual Arts, MPS Branding Program, New York, NY

Third System™ (research entity)

Methodology

This report presents findings from Third System's AI Availability research program. The AI Availability Score (AIAS) is an early-stage measurement framework introduced in Tri-System Brand Growth (Gonzalez Castro, 2026). Findings should be used directionally. Construct validity — the correlation between AI Availability and consumer behavior — has not yet been established and is the subject of Phase 3 of the research program. Methodology is published in full and versioned at thirdsystem.ai/methodology (current: v1.7). Registry construction is curated and reflects the framework's view of each category's competitive set; revisions are documented in the methodology log.

Citation

González Castro, P. U. (2026). *Consistency, Category by Category: A Cross-Category Baseline for the AIAS Consistency Component (CPC), v0.31*. Third System (SSRN pending). thirdsystem.ai/v31

Underlying datasets

Per-brand CPC: osf.io/ec6wh/v31/data/v31_cpc.csv

Verdicts: osf.io/ec6wh/v31/v31_cpc_verdicts.json

Scoring code: osf.io/ec6wh/v31/scripts/score_v0_31.py

Methodology log: v0.31 (SSRN pending).

Inquiries: hello@thirdsystem.ai