

When the Models Change, What Holds?

AI brand-consistency rankings tested across a major model upgrade: the broad order holds, the score itself does not reproduce, and the instability concentrates in whichever model family jumps furthest — with consequences for tracking a brand over time.

June 2026 · Pablo Ulpiano González Castro

Third System™ · AIAS™ Measurement Program · v0.32

Generative AI increasingly decides which brands surface when a buyer asks for a recommendation — and the models doing the deciding are retired and replaced every few months. That puts a blunt question ahead of any AI-brand metric: when the model changes, does the reading change with it? We tested one consistency measure across a deliberately large model upgrade. The answer is useful, with a caution attached.

An AI-availability reading earns its keep only if it tracks the brand, not the tool beneath it. To test whether one does, we held twenty-four automotive brands fixed and measured their AI-presence consistency twice — once on an older generation of the six-model panel we use, once on the current frontier — changing nothing but the models. The jump was deliberately large, to stress the measure rather than flatter it. Three findings came back, and a fourth to watch. The broad order of brands held across the upgrade, but only just. The consistency score for any given brand did not hold: it moved, on average, by more than the gap separating one brand from the next. The instability was not evenly spread — it concentrated

in the single model family that made the largest version jump. And the newest electric challengers are beginning to register in AI recall, though not yet enough to move the measure.

Standings broadly survive an upgrade — weakly. The order of brands on AI consistency holds across a model change, but the margin is thin. Read a single ranking as indicative, not settled. (P1)

The score does not reproduce. A brand's consistency reading is not stable across a model change to within the measure's own resolution — it can move more than the spacing between brands. Don't over-read a period-to-period score change. (P2)

The wobble has an address. Instability concentrates in whichever model family makes the largest version jump; a panel is only as stable as its most volatile member. (P3)

Emerging brands: watch, don't act. Electric challengers are gaining AI recall under newer models but remain below the threshold this measure reads — a movement to watch, not yet

a measured effect. (P4)

Comparability is a choice. A consistency reading is comparable across periods only on a version-controlled panel — pinned, dated models migrated deliberately — or the score partly reflects the tools rather than the brands. (P5)

What we measured

We measured consistency: how steadily a brand surfaces across the different AI models a buyer might use, rather than depending on which model happens to answer. The test bed was twenty-four automotive brands — a deliberate mix of heritage names, electric challengers, mass-market makes, and discontinued marques — read across a panel of six models from three providers. We ran the full measurement twice: once on an older generation of those six models, once

on the current generation, holding the brands, the questions, and everything else fixed. Only the models changed, and we made the gap between old and new deliberately large so the test would expose version sensitivity rather than hide it. One note on what “consistency” counts: it reads from how often each model names a brand unprompted across a battery of category and cultural questions; a brand the models rarely name falls below a floor and carries no reading at all.


FINDING 01

Standings hold — but only weakly

Standings hold their broad order across a model upgrade — but weakly

Per-brand AI-consistency: older model panel vs current

Treat a single reading as indicative, not definitive — the order shifts with the panel.



Proposition 1 — standings hold across an upgrade, but weakly.

Source: AIAS™ v0.32 | Protocol v1.7 | Third System™

P1 · Standings hold, weakly. Brands ranked by AI-consistency on the older model panel against the current one. The broad order holds (rank agreement 0.71), but it is weak and depends on which model family is in the panel — Finding 3 decomposes that dependence.

Ranked by consistency, the brands sit in nearly the same order on the old and new models (rank correlation 0.71). But the margin is slim and the result leans on two of the three providers; the third's models, taken alone, would put the orders in substantial disagreement.


FINDING 02

The score does not reproduce

A consistency score is not reproducible across an upgrade

Per-brand score change, older panel to current

The typical shift is larger than the gap between brands — so period-over-period changes must not be over-read.



Proposition 2 — score not reproducible to the measure's own resolution.

Source: AIAS™ v0.32 | Protocol v1.7 | Third System™

P2 · The score does not reproduce. Per-brand change in consistency from the older panel to the current. The typical move is larger than the gap separating one brand from the next, with no overall direction — newer models reshuffle which brands read as consistent rather than shifting the level.

Across the brands that carried a reading, the average score moved by 0.077 between old and new models — larger than 0.069, the mark for half the spread between brands. The


moves had no net direction (they averaged near zero): newer models reshuffled which brands read as consistent rather than raising or lowering all of them.

FINDING 03

The wobble has an address

The wobble has an address

Leave-one-model-family-out sensitivity of the standings agreement (Finding 1)
Drop the family that made the largest version jump and the agreement clears the line; drop a steadier family and it falls below.



Proposition 3 — a robustness check on Finding 1's standings agreement, not an independent measurement.
Source: AIAS™ v0.32 | Protocol v1.7 | Third System™

P3 · The wobble has an address. Leave-one-model-family-out sensitivity of the standings agreement — a robustness check on Finding 1, not an independent result. Dropping the family that made the largest version jump lifts agreement above the line; dropping a steadier family drops it below.

Removing the provider whose models jumped furthest — an older-to-current leap spanning a full model-family change — raised the ranking agreement from 0.71 to 0.82; removing either other provider dropped it to about 0.5. The instability lives in the largest jump.


FINDING 04

Emerging brands: watch, don't act

Emerging brands are gaining AI recall — a movement to watch

New-brand recall, older panel vs current

Rising under newer models, but below the level this measure registers: watch it, do not yet act on it.



Proposition 4 — a signal to watch, not a measured effect.

Source: AIAS™ v0.32 | Protocol v1.7 | Third System™

P4 · A movement to watch, not measured. Emerging-brand AI recall, older panel versus current. No brand crossed the threshold the measure reads (top); recall is rising beneath it (bottom). A signal to watch — not yet a measured effect.

Rivian and Lucid were recalled essentially never by the older models and began to be named by the newer ones — Rivian's average mentions rose from zero to roughly two-thirds of one

across the panel. Real movement, but below the threshold at which the measure begins to read, so it registers as no change.

Scope and boundaries

Because the score moves with the models (P2) and the movement tracks the largest version jump (P3), a reading taken on one panel vintage is not directly comparable to one

taken on another. Comparability is recoverable only by fixing the panel's versions and migrating them deliberately.

Four patterns

Four patterns came back. First, the standings broadly held: rank the brands by consistency on the old models and on the new, and the two orders largely agree — but the agreement is thin, and fragile. Second, the score itself did not hold. A given brand's consistency reading moved between old and new models by more than the typical gap separating one brand from the next, and in no consistent direction: newer models did not make brands look uniformly more or less consistent, they reshuffled which brands read as consistent. The measure is reproducible enough to rank by once, but not to track period

over period. Third, the instability has an address — it concentrated in the single provider whose models made the largest version jump; remove that provider and the standings snap back into agreement. A panel is only as steady as its most volatile member. Fourth, and to watch rather than bank: the electric challengers are rising. Rivian and Lucid, recalled essentially never by the older models, began to surface under the newer ones — but not often enough to cross the floor where the measure starts reading. The direction is right; the magnitude isn't there yet.

Limitations

Read these within their bounds. This is one category — automotive, chosen for its spread of brand types; consistency may behave differently where the surrounding discourse is denser or thinner. It is one moment in time: we changed the models but not the calendar, so this isolates the effect of upgrading the models, not drift that would occur anyway. And it rests on the fourteen brands that carried a reading; the other ten — discontinued marques, the newest EVs, and even

a mainstream name like Nissan — sat below the recall floor, a reminder that this measure tracks how often a brand is named, not how large it is. Most important to weigh: the consistency measure itself is an early version, one we are actively redefining. These results are partly a stress test of the gauge and feed that redefinition as much as they describe the market. Treat the propositions as directionally sound and the exact scores as provisional.

What's next

Three moves follow. We are redefining the consistency measure to remove a known dependence on how often a brand is recalled, and the next step is to re-run this exact old-versus-new test on the redefined version — version-robustness becomes a bar any future gauge must clear before it can be trusted over time. We will widen the test beyond automotive, to see whether the partial, provider-dependent stability here holds in categories with denser or sparser discourse. And we will keep watching the

challengers below the floor: a sub-threshold rise, tracked deliberately, is the kind of signal that becomes a measured shift before competitors notice it. For anyone using an AI-availability reading operationally, the immediate move is the fifth proposition — control the panel: pin the exact model versions you measure on, migrate them on a deliberate schedule rather than letting them update underneath you, and record which versions produced each reading. Otherwise a change in a brand's score is partly a change in the instrument.

The five propositions, scored

Each proposition is the managerial reading of a locked result, scored for brand teams rather than reviewers.

#	Proposition	Status
P1	Standings hold across a model upgrade	Holds — weakly
P2	The score reproduces across a model upgrade	Does not hold
P3	Instability is spread across the panel	No — it concentrates in the largest jump
P4	Emerging brands register on the measure	Not yet — watch
P5	A reading is comparable over time as-is	No — comparability needs a version-controlled panel

The headline is not that AI brand-consistency is stable across model upgrades, nor that it collapses. It is that the standings survive while the scores do not — and the difference governs how the reading should be used. As a one-time ranking, the measure separates brands meaningfully. As a tracker, it cannot yet tell a real change in a brand from the turnover of the models beneath it, unless the panel is held still. For now, that is the operative instruction: pin the versions, migrate them on your terms, and record what produced each reading. The gauge is being sharpened; until it is, treat the order as signal and the scores as provisional — and don't mistake the instrument moving for the market moving.

About the Third System

The Third System describes a shift in how consumers make decisions: instead of browsing and comparing options directly, many now rely on AI systems to generate recommendations. In this environment, brands compete not only for human attention but for inclusion within AI-generated outputs. This creates a new competitive layer where visibility is determined by how AI systems interpret and surface brands.

Authored by

Pablo Ulpiano González Castro

School of Visual Arts, MPS Branding Program, New York, NY

Third System™ (research entity)

Methodology

This report presents findings from Third System's AI Availability research program. The AI Availability Score (AIAS) is an early-stage measurement framework introduced in Tri-System Brand Growth (Gonzalez Castro, 2026). Findings should be used directionally. Construct validity — the correlation between AI Availability and consumer behavior — has not yet been established and is the subject of Phase 3 of the research program. Methodology is published in full and versioned at thirdsystem.ai/methodology (current: v1.7). Registry construction is curated and reflects the framework's view of each category's competitive set; revisions are documented in the methodology log.

Citation

González Castro, P. U. (2026). *Consistency, Category by Category: A Cross-Category Baseline for the AIAS Consistency Component (CPC), v0.32*. Third System (SSRN pending). thirdsystem.ai/v32

Underlying datasets

Per-brand CPC: osf.io/ec6wh/v32/data/v32_cpc.csv

Verdicts: osf.io/ec6wh/v32/v32_cpc_verdicts.json

Scoring code: osf.io/ec6wh/v32/scripts/score_v0_32.py

Methodology log: v0.32 (SSRN pending).

Inquiries: hello@thirdsystem.ai