

The Provider Fingerprint

Does the company that built an AI model leave a systematic mark on how consistently it recalls brands? A pre-registered re-analysis across five categories finds a real but faint signal - and a measurement blind spot worth naming.

June 2026 · Pablo Ulpiano González Castro

AIAS™ Measurement Program

Three companies build the large language models most brands are now seen through. We asked whether that authorship shows up as a consistent pattern in how brands are recalled – and found that the more useful answer was about the question we could not yet ask.

AI Availability – how a brand surfaces inside large language models – is the third front of brand availability, alongside the mental and physical availability marketers already manage. Within it, Consistency asks how steadily a brand is recalled across the models people actually use. Holding the models fixed at a single moment, this study asked a narrower question: does the provider behind a model – Anthropic, OpenAI, or Google – leave a systematic fingerprint on that consistency?

The short answer is yes, but faintly – and the more important answer is that we could not test the question that matters most.

Across 112 brands in five categories, which provider built a model does explain a real, repeatable share of the variation in how consistently brands are recalled. The effect is reliable but small: provider sits just above what the structure of the test would produce by chance, and no further.

We then asked the sharper question – whether that provider pattern reflects consistency over and above simple recognition, or merely restates which brands the models already know. We could not answer it. Among the brands a model recalls at all, recognition is effectively universal – every model recognizes them – so there is no provider pattern in recognition to compare against. The test we had pre-registered to settle this passed its arithmetic, but the comparison was empty. Reporting that honestly, rather than as a win, is the study's most useful result.

No single provider proved reliably the most consistent across categories, and the smaller differences seen for lesser-known brands turned out to be mechanical rather than meaningful.

What we measured

This is a re-analysis: no new model queries were run. It re-examines frozen measurements from an earlier cross-category study - 112 brands across audiophile headphones, skincare, cosmetics, automotive, and premium spirits - scored across a fixed panel of six models, two each from Anthropic, OpenAI, and Google. For each brand we asked how much of the model-to-model variation in recall is

organized by provider rather than scattered within each provider's own pair, and whether the same pattern appears in plain recognition. The full method, the questions to be tested, and the rules for judging them were locked and time-stamped before any scoring, and every figure was checked back against the original study.

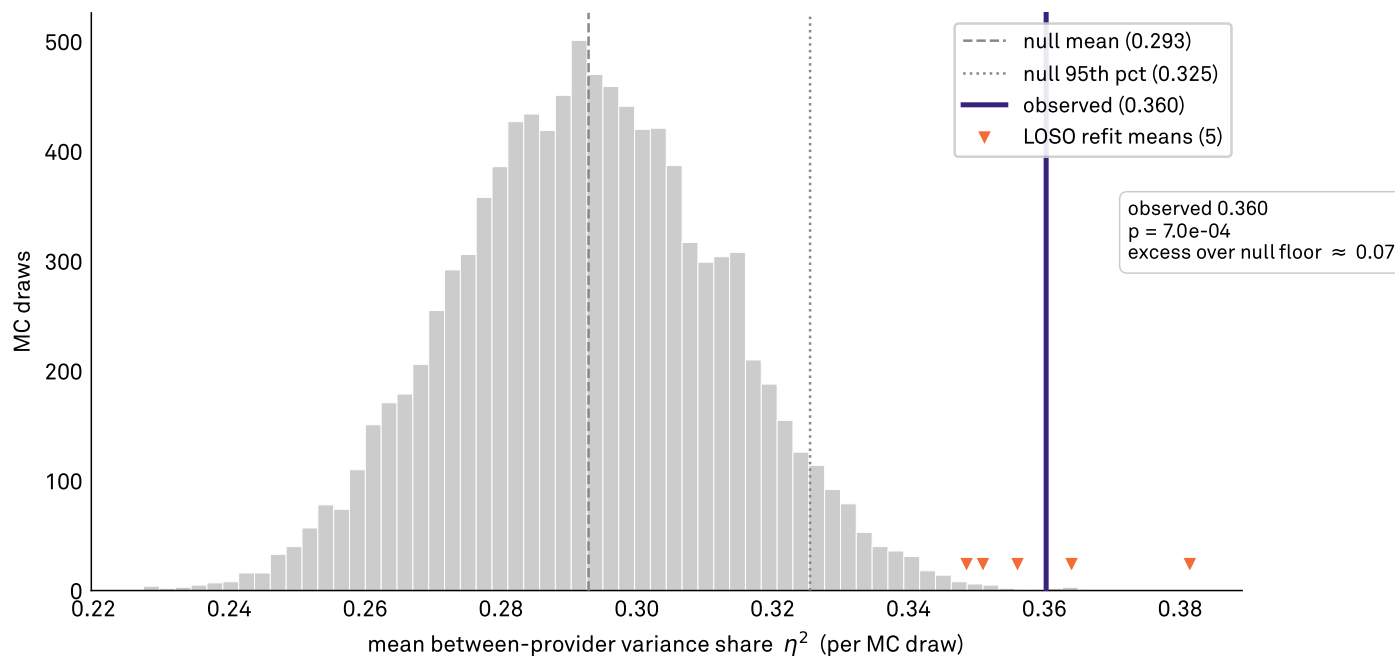
FINDING 01

A real fingerprint, faintly pressed

Provider organization of recall is significant but modest

Permutation-null of mean η^2 across 112 brands ($\geq 10,000$ draws); observed vs the null

The observed share sits just past the 95th percentile of a high chance floor — a reliable but small provider signal.



P1 SUPPORTED — *provider shapes recall consistency reliably but modestly; a small excess over a high chance floor.*

Source: AIAS™ v0.33 (Provider-Asymmetric CPC) | Protocol v0.33-prereg-r1 | Third System™

A real fingerprint, faintly pressed. The provider share of recall variation sits just past the 95th percentile of its chance distribution — reliably above chance across all five categories, but only modestly.

The provider share of recall variation is reliably above chance and holds across all five categories, surviving every leave-one-category-out recheck. Its size is small - just above

the high floor the three-pair structure imposes - so it reads as a dependable but minor influence.

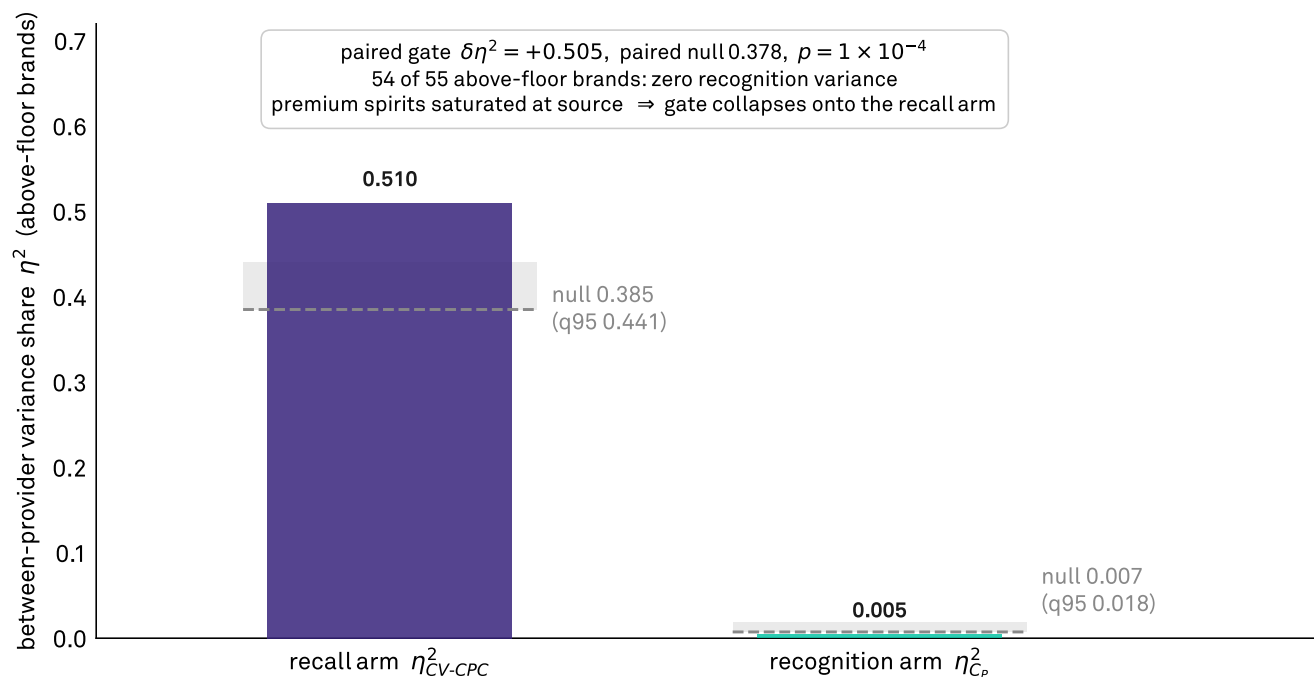
FINDING 02

The question that had no answer

The Beyond-Presence gate collapses under recognition saturation

Recall-arm vs recognition-arm η^2 among above-floor brands, each with its permutation null

Recognition carries almost no provider variance ($\eta_{C_p}^2 \approx 0$), so the paired gate measures recall asymmetry, not a contrast.



P2 NOT ESTABLISHED — the beyond-recognition test could not be run; recognition is saturated among recalled brands, so the gate measures recall alone.

Source: AIAS™ v0.33 (Provider-Asymmetric CPC) | Protocol v0.33-prereg-r1 | Third System™

The question that had no answer. The recall arm carries provider structure; the recognition arm carries almost none, because every model recognizes these brands — so with nothing to compare against, the gate measures recall alone.

The pre-registered gate met its number, but only because recognition contributed essentially nothing to compare against: among recalled brands, every model recognizes the brand, leaving no provider pattern in recognition. The gate

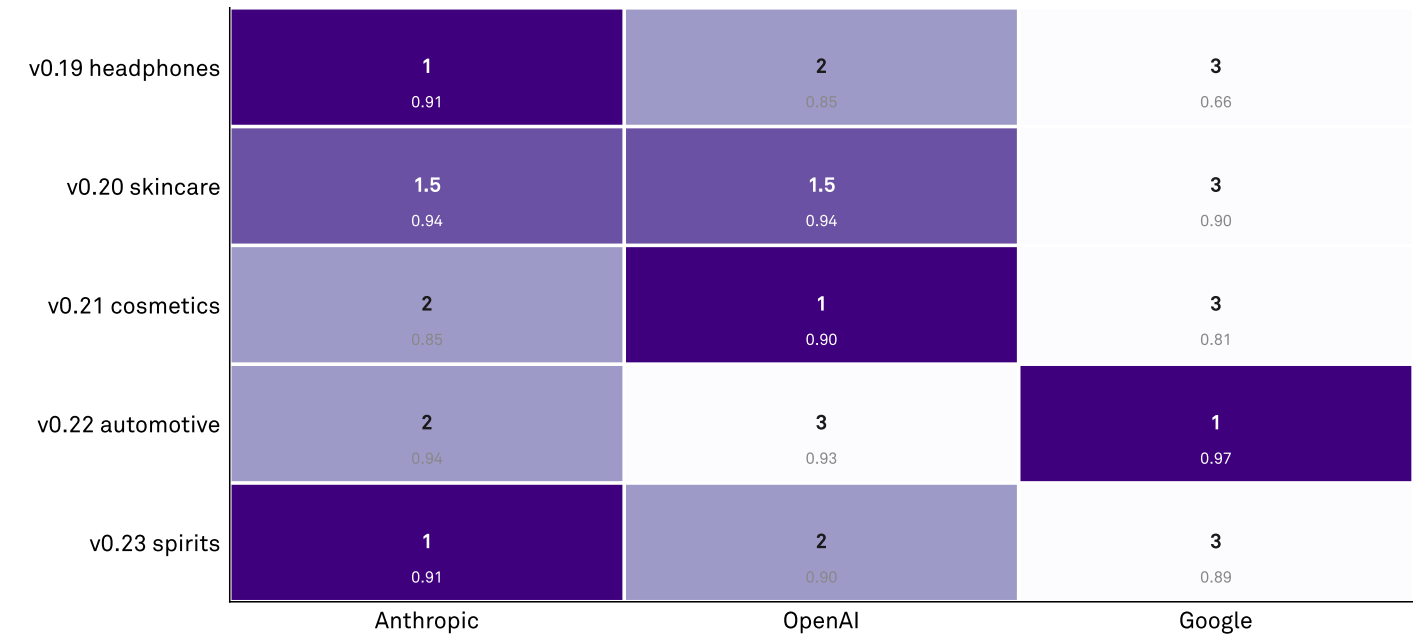
therefore measured recall asymmetry alone, not a separation from recognition. We treat the beyond-recognition question as unanswered.

FINDING 03

No provider stands out as steadiest

Provider consistency ranks are not stably ordered

Within-pair-consistency provider rank by substrate (1 = most consistent; cell sub-value = consistency)
Kendall's W = 0.31 against the exact null ($p = 0.18$) — underpowered at five substrates, uninformative about concordance.



P3 NOT SUPPORTED — no provider is reliably steadiest; rankings reshuffle and five categories cannot resolve them.
Source: AIAS™ v0.33 (Provider-Asymmetric CPC) | Protocol v0.33-prereg-r1 | Third System™

No provider stands out as steadiest. Consistency ranks reshuffle across categories (skincare is a tie at the top); with only five categories, the pattern cannot be told from noise.
Provider consistency rankings change from category to category, and five categories are too few to distinguish a true ordering from noise. The result is read as uninformative rather than as evidence of parity.

FINDING 04

Obscurity is not a signature

Barely-recalled brands show smaller provider differences, but this follows mechanically from their having little recall variation to organize. It is not a distinct property of obscure brands.

Read the lens before the brand is known

Taken together, the provider lens is meaningful only where recognition still varies - the emerging-recognition regime. For universally recognized brands, today's recognition signal is too

saturated to support the comparison. This is a scoping rule for measuring AI Availability Consistency, and the study's most portable conclusion.

In summary

Three things to take away. Provider authorship shapes how consistently models recall a brand - reliably, but faintly. Whether that shaping is about consistency rather than familiarity could not be tested, because recognition is already

universal among the brands models recall. And no single provider proved the steadiest. The faint and the unanswerable, not a bold provider effect, are the honest results.

Limitations

Three cautions frame these findings. First, scope: five categories and a six-model panel can detect whether a provider effect exists, but not rank the providers with confidence or generalize broadly. Second, the consistency score used here is an inherited, not-yet-validated measure, treated as-is; this study does not certify it. Third, and most consequentially, the recognition signal available to us is

all-or-nothing and saturated among recalled brands, so the 'beyond recognition' question is bounded by the instrument rather than settled - a graded recognition signal simply does not exist in this data to push further. The provider effect is also small, and should be read against the high baseline the test's structure imposes.

What's next

The clearest next move is a recognition signal with room to vary. The beyond-recognition question is answerable only where recognition is not already saturated, so a graded or scaled recognition measure - which this data does not contain - is the prerequisite for any future claim that provider shapes

consistency rather than mere familiarity. Beyond that, a wider set of categories would support the provider ranking this study could not resolve, and more models per provider would lift the test off its high chance floor. Each feeds the program's next-generation consistency instrument.

The five propositions, scored

Each proposition is the managerial reading of a locked result, scored for brand teams rather than reviewers.

#	Proposition	Status
P1	Provider identity shapes how consistently models recall a brand.	SUPPORTED — real but modest
P2	That provider effect reflects consistency beyond simple recognition.	<i>NOT ESTABLISHED — untestable here - recognition saturated</i>
P3	One provider is reliably the most internally consistent across categories.	<i>NOT SUPPORTED — unresolved; underpowered at five categories</i>
P4	Lesser-known brands carry a distinct provider signature.	<i>NOT SUPPORTED — the pattern is mechanical, not a signature</i>
P5	The provider lens is informative only where recognition is not yet saturated.	SUPPORTED — the study's central lesson

The most valuable findings sometimes map their own edges. Provider authorship leaves a faint, real mark on how consistently AI models recall a brand - worth knowing, easy to overstate. But the question brands will most want answered - is this about consistency, or just familiarity? - turns out to be unanswerable for the brands that need it least, and answerable only for those still earning recognition. That boundary, not the faint signal, is what to carry forward: read the provider lens where recognition is still being won, and treat it with caution once it has been.

About the Third System

The Third System describes a shift in how consumers make decisions: instead of browsing and comparing options directly, many now rely on AI systems to generate recommendations. In this environment, brands compete not only for human attention but for inclusion within AI-generated outputs. This creates a new competitive layer where visibility is determined by how AI systems interpret and surface brands.

Authored by

Pablo Ulpiano González Castro

School of Visual Arts, MPS Branding Program, New York, NY

Third System™ (research entity)

Methodology

This report presents findings from Third System's AI Availability research program. The AI Availability Score (AIAS) is an early-stage measurement framework introduced in Tri-System Brand Growth (Gonzalez Castro, 2026). Findings should be used directionally. Construct validity — the correlation between AI Availability and consumer behavior — has not yet been established and is the subject of Phase 3 of the research program. Methodology is published in full and versioned at thirdsystem.ai/methodology (current: v1.7). Registry construction is curated and reflects the framework's view of each category's competitive set; revisions are documented in the methodology log.

Citation

González Castro, P. U. (2026). *Consistency, Category by Category: A Cross-Category Baseline for the AIAS Consistency Component (CPC), v0.33*. Third System (SSRN pending). thirdsystem.ai/v33

Underlying datasets

Per-brand CPC: osf.io/ec6wh/v33/data/v33_cpc.csv

Verdicts: osf.io/ec6wh/v33/v33_cpc_verdicts.json

Scoring code: osf.io/ec6wh/v33/scripts/score_v0_33.py

Methodology log: v0.33 (SSRN pending).

Inquiries: hello@thirdsystem.ai